

# Road traffic congestion detection and tracking with Spark Streaming Analytics

Thorsteinn Thorri Sigurdsson ([ttsi@kth.se](mailto:ttsi@kth.se))

Supervisors: Ahmad Al-Shishtawy & Zainab Abbas  
Examiner: Vladimir Vlassov

# Motivation and goal

- Traffic congestion poses several problems
  - Increased pollution and fuel consumption
  - Increased commuting times with higher transportation costs
  - Reduced traffic safety
- Goal
  - Real-time detection and tracking of congestion in road system
  - Using data from infrastructure traffic sensors



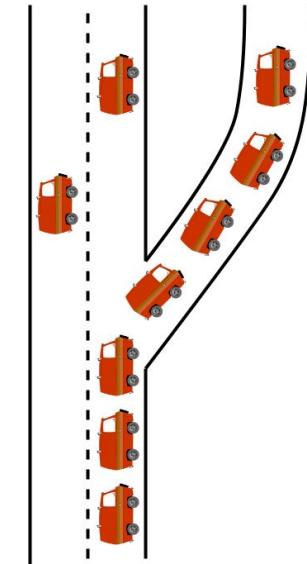
# Benefits

## Real time congestion detection use cases

- Send warnings to drivers approaching end-of-queue
- Congestion mitigation
  - Adaptive speed limits
  - Re-routing of traffic

## Congestion tracking use cases

- Queue spill-back analysis
- Optimization of road network by traffic authorities



# The data set

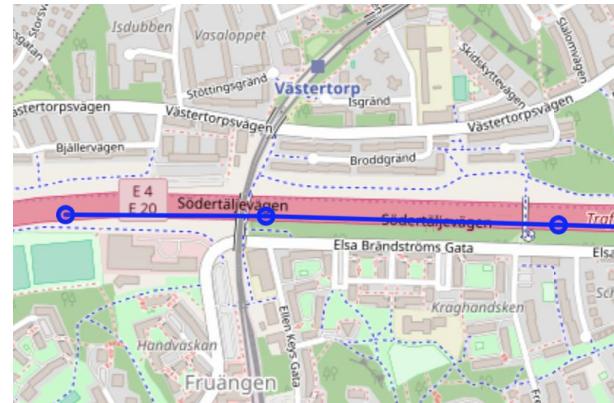
- 2059 radar sensors placed on main arteries in Stockholm (+ a single road in Göteborg)
- 150-400 meters apart
- Sensor measurement every minute
  - Flow (vehicles/minute)
  - Average speed (km/h)
  - -> Calculated density (vehicles/km)
- Sensors identified by
  - Road ID
  - Relative kilometer reference
  - Lane ID (sensor for each lane on the road)
  - + GPS coordinates for each sensor



*E4N 55650. Lanes 1-4.  
Image from Google Maps' Streetview.*

# Problem statement

- Instead of detecting congestion at individual sensors, extend detection into spatial dimension
- Detect regions of congestion
  - Connected congested sensors representing a queue
- Track evolution of queues through time
  - Build up, split, movement, dissipation
- Resolution down to the individual lane level



# Problem statement

- Construct traffic sensor graph
- Frame the congestion detection problem as a graph processing problem
- Adapt existing graph processing algorithms for this task
  - Graph clustering
  - Community detection
- Use streaming graph processing algorithms for real time congestion detection and tracking

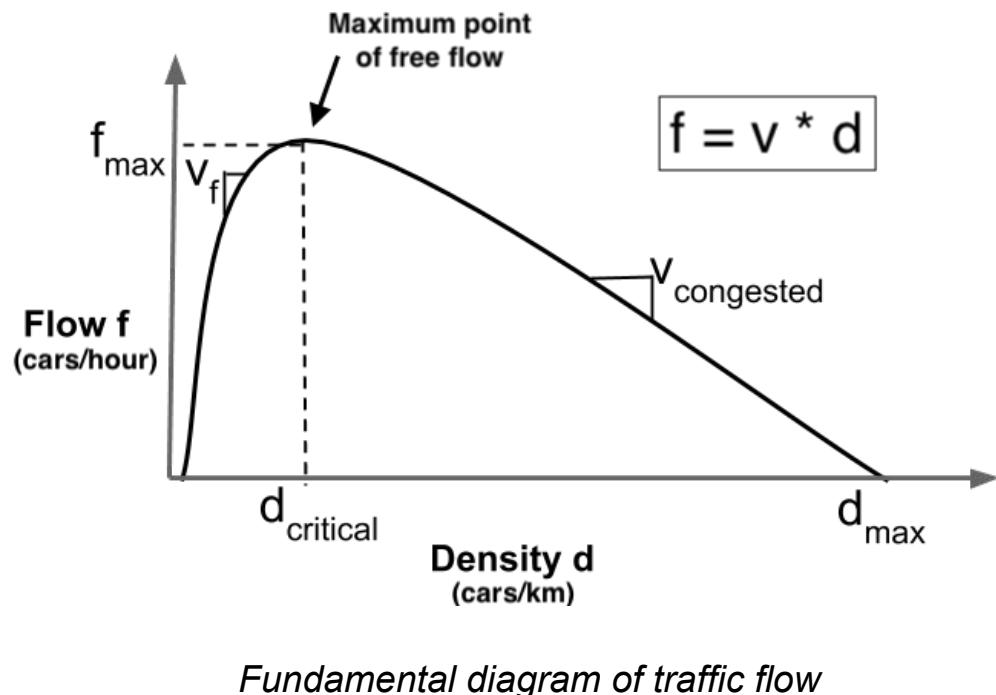
# Background - Traffic theory

3 main metrics

- Flow (vehicles / hour)
- Speed (km / hour)
- Density (vehicles / km)

Traffic flow can be classified as:

- Free flow
  - Vehicle can do any maneuver the driver wishes
- Congestion
  - Complement of free flow



# Background - Traffic queues

- Can be defined as a row of vehicles waiting to be served
  - Moving at slow speeds or stopped
- Safety hazard for cars approaching the end-of-queue at higher speeds
- Can spread to disrupt traffic on other roads otherwise un-congested
  - Queue spillback
- We take a connected chain of congested traffic sensors to represent a queue



# Background - shockwaves

A shockwave is a change or discontinuity in traffic conditions

- Propagation of change in speed or density
- Travels through the road system

Queue buildup at an intersection is a shockwave (queue extends upstream)

Can also form “out of thin air” when density exceeds critical threshold



*Sugiyama et al.*

# Related work

- Traffic congestion can be detected with a number of different approaches
- Microscopic methods
  - Can view traffic flow at individual vehicle basis
    - Probe vehicles
    - Vehicle re-identification methods
    - Vehicle to vehicle communication systems
- Macroscopic methods
  - View traffic flow in aggregate
    - Cameras and image recognition
    - Acoustic sensors
    - Aerial and satellite imagery
    - Aggregated traffic statistics (average speed, flow, density)

# Related work

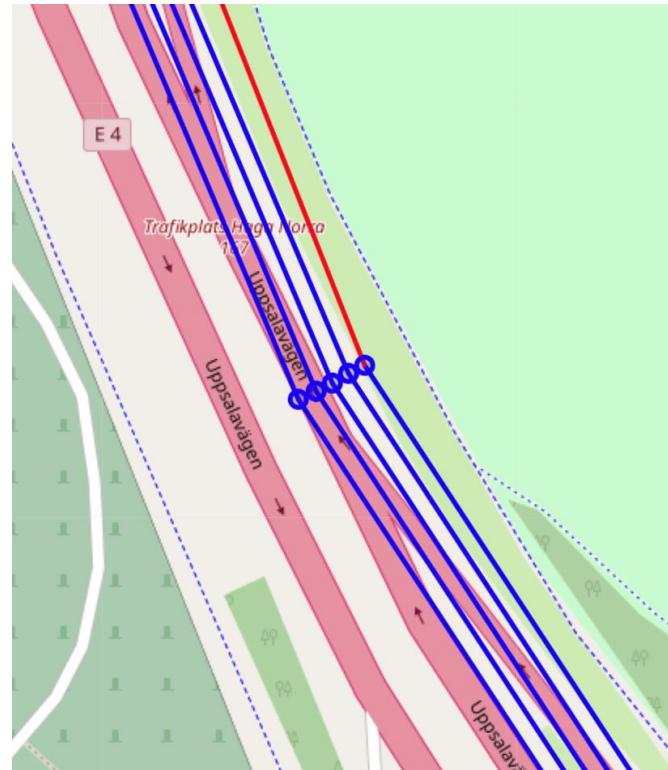
- Google Maps
  - Probe vehicles
  - Speed measurements from smartphones collected
  - Red, yellow, green roads in Google Maps application
- Coifman (2003)
  - Congestion detection with infrastructure sensors
    - Dual loop induction sensors
    - Give vehicle length measurements
  - Re-identification of vehicles at different locations, by vehicle length
  - Abnormally high link journey times -> congestion
  - Ground truth: Video recordings

# Related work

- Li et al. (2007)
  - FlowScan algorithm
    - Density based clustering to find “hot routes”
    - i.e. not congestion detection
  - Probe vehicle trajectories
  - Directed graph
    - Vertices are intersections or landmarks
  - Uses number of shared trajectories and number of hops between edges for density based clustering
  - Ground truth
    - Traffic simulator data
  - Batch processing only

# The traffic sensor graph

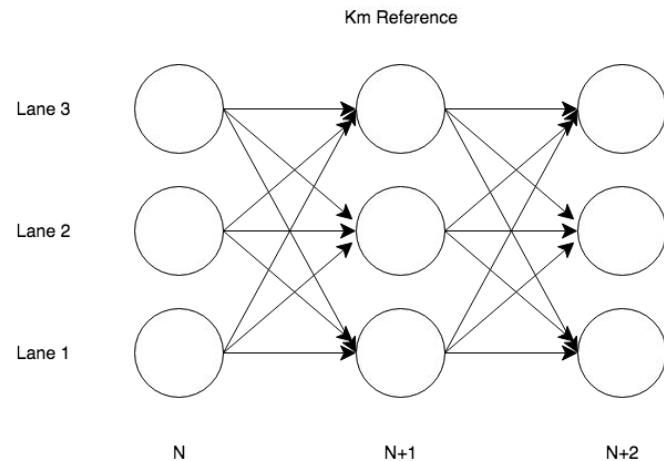
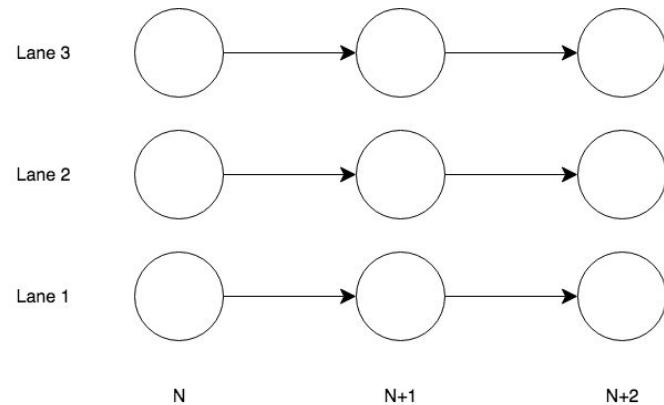
- Directed graph modelling spatial relationship between traffic sensors
- Sensors as vertices, road segments between sensors as edges
- A path in the graph is a possible path to drive in road system
- At most 2037 sensors
  - Graph valid from 2016-5-9 onwards



*E4N next to Frösunda. Red edge is a connection to another road, E4A.*

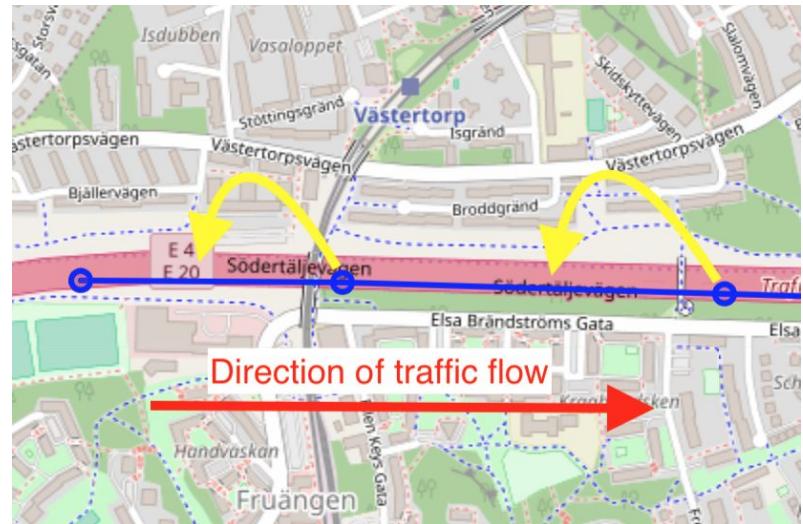
# The traffic sensor graph

- Two types of graphs were constructed
- Base graph
  - No lane changes
  - Used in congestion detection and tracking, per lane
- Reachability graph
  - Allow lane changes
  - Can be used to send “congestion ahead” warnings to drivers

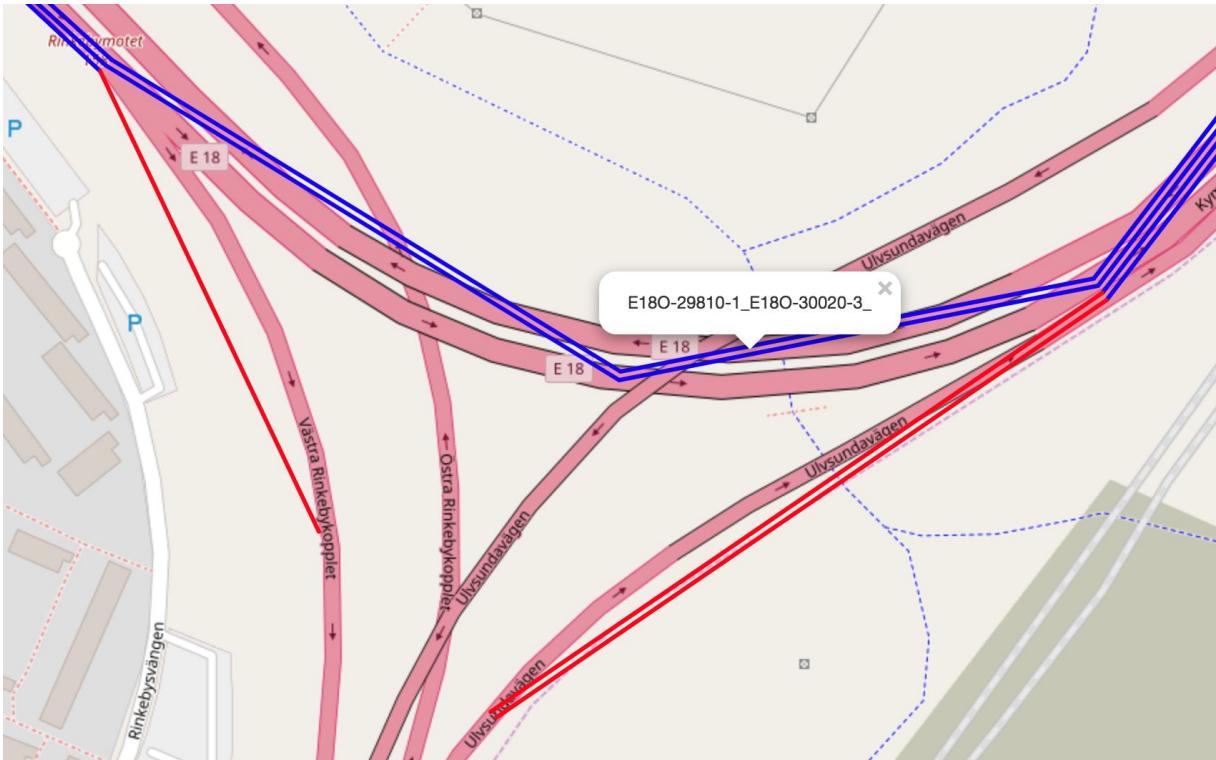


# The traffic sensor graph

- Directed graph
  - Each edge has a source and destination vertex
- Edges weighted with sensor measurements from destination vertex
  - Assumed to represent the conditions on the road segment upstream of the destination vertex
- Weights updated every minute (new graph)

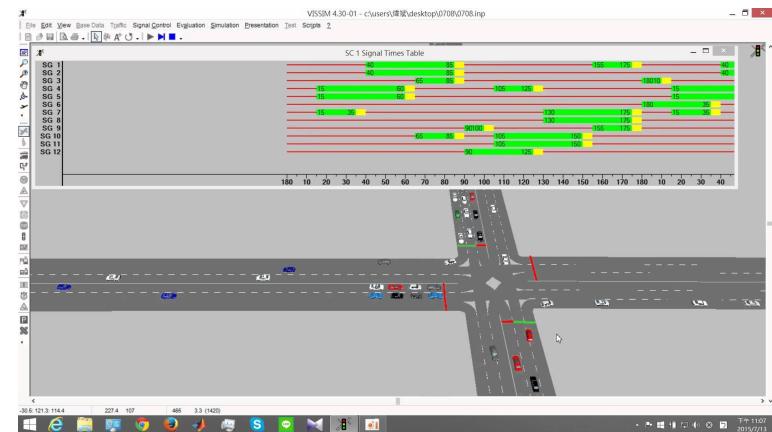


# The traffic sensor graph



# Ground truth

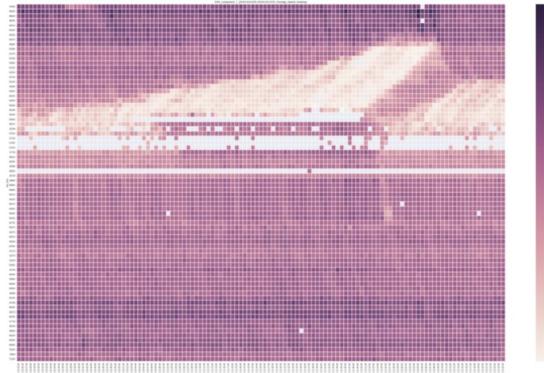
- All implemented methods unsupervised
- No ground truth available for our data set
  - What constitutes congestion?
- Related work has used
  - Video recordings of traffic
  - Domain experts to identify and label ground truth
  - Fusion with external data sets such as accidents
  - Traffic simulators
    - Allowing complete observation of synthetic traffic flow
    - Vehicle trajectories, number of vehicles in queue etc.



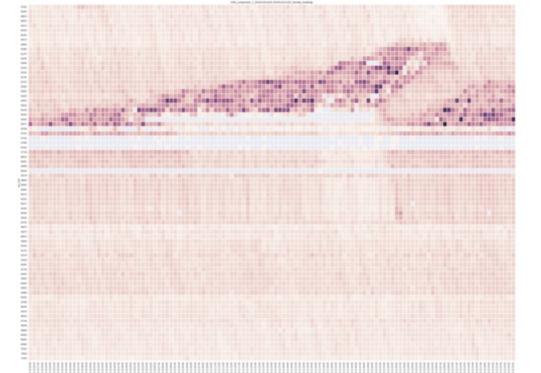
*Screenshot of the VISSIM Traffic Simulator*

# Ground truth

- Use visual inspection of heat maps to compare results to underlying data
- Spatio-temporal patterns of measured traffic features
  - Time on x-axis, kilometer reference on y-axis (space)
- Congestion detection methods should find the same patterns
- Domain expert approach



*Average speed*



*Density*

# Manually labelling data set

Find connected components in base graph

- Represent “actual lane” - path if driving without changing lanes
- Can’t rely on lane IDs of sensors

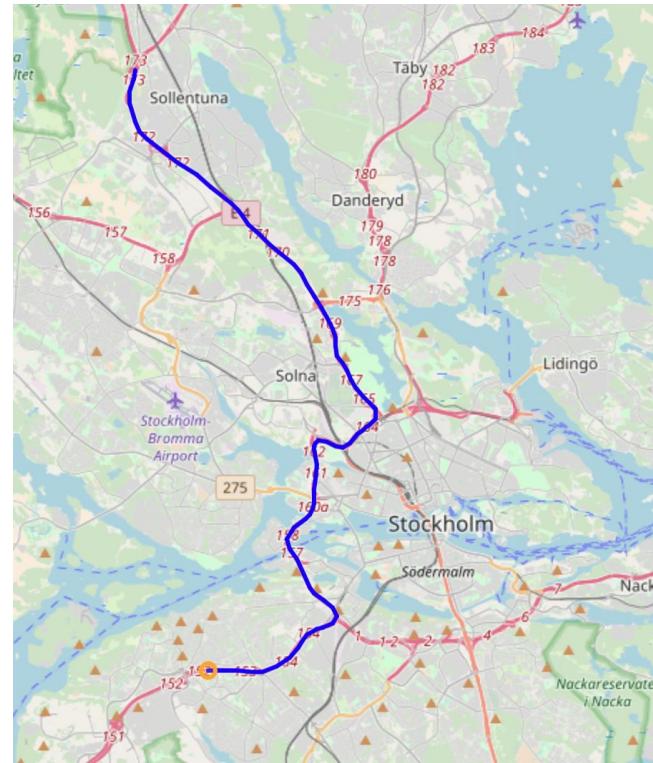
8 test congestion patterns identified on the “actual lanes”

- Queues and shockwaves labeled
- 16 queues, 15 shockwaves

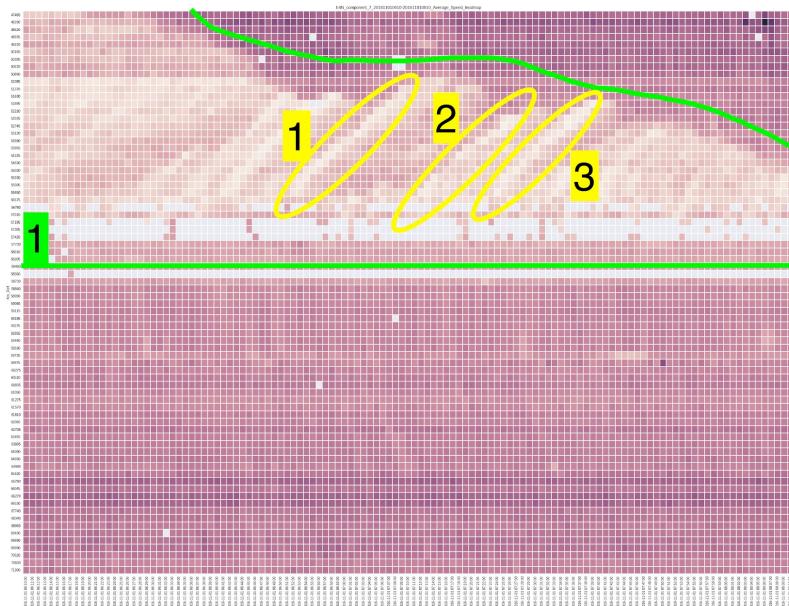
Criteria for selection

- Combined they should cover most of the road network
- Show interesting congestion patterns

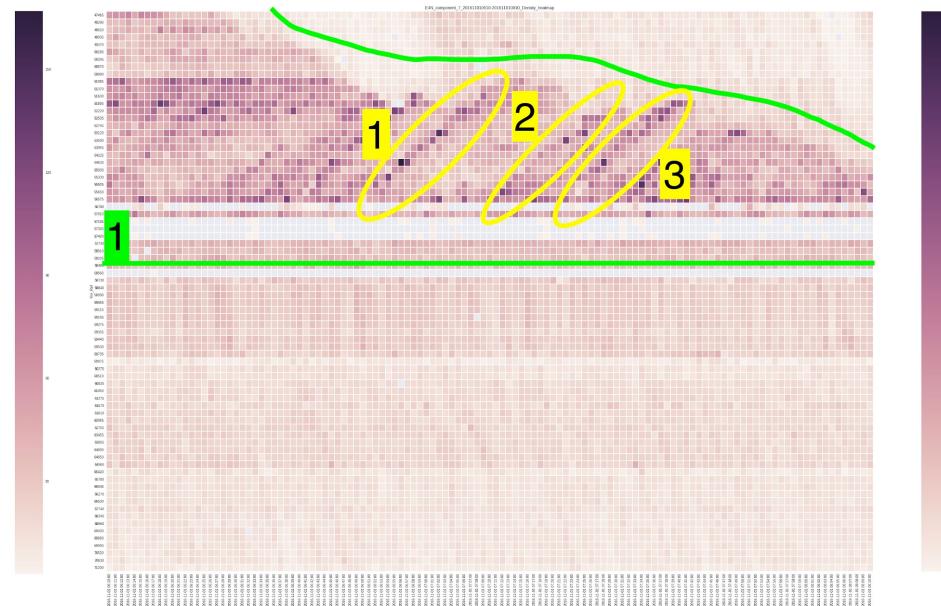
# E4N 2016-11-1 (component 7)



# E4N 2016-11-1 06:10 - 08:10 (component 7)

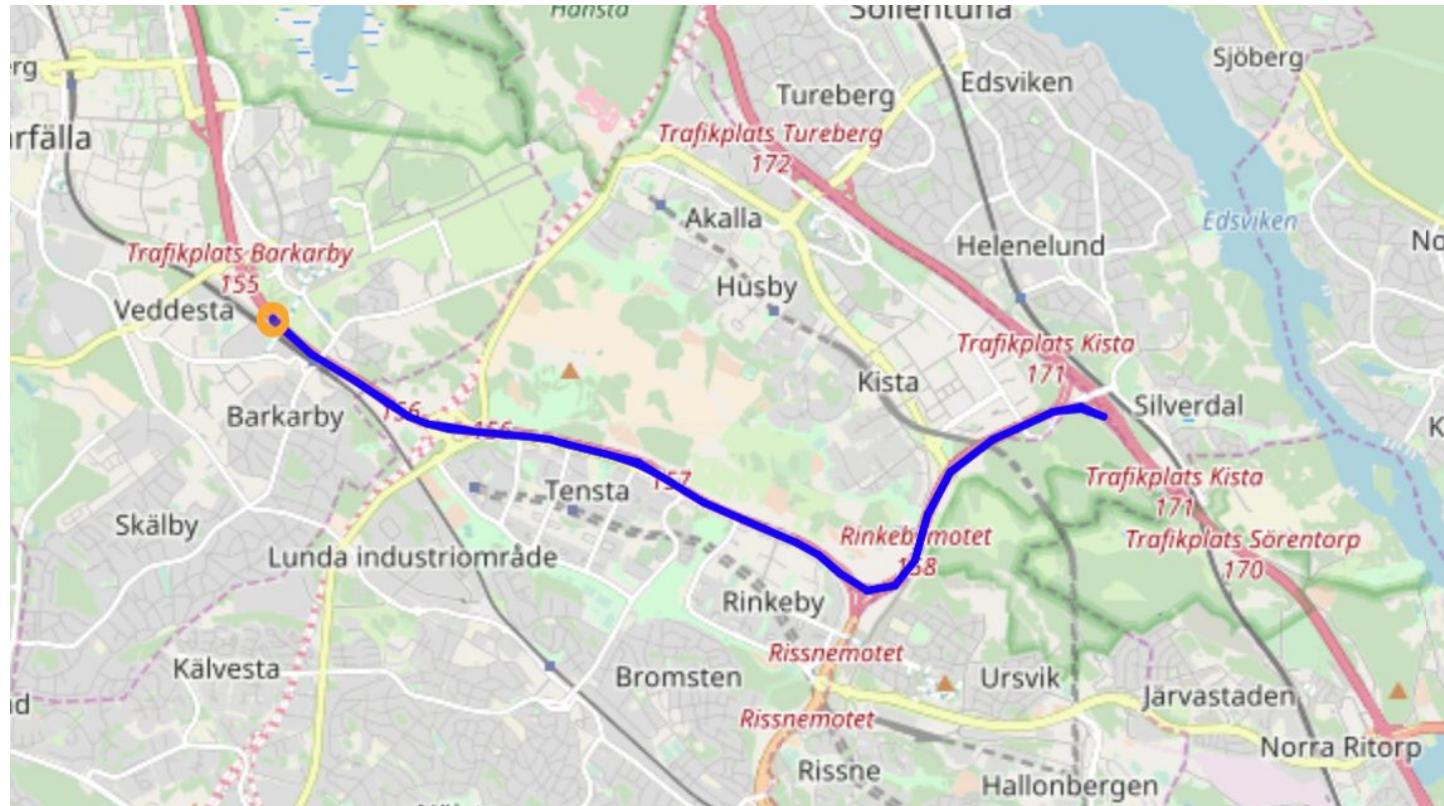


*Average speed*

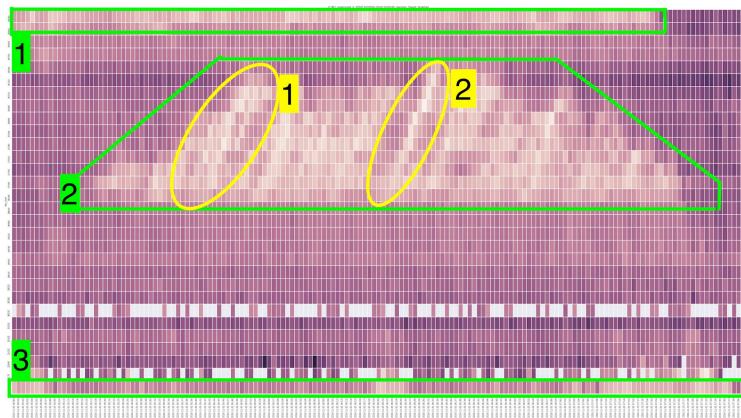


*Density*

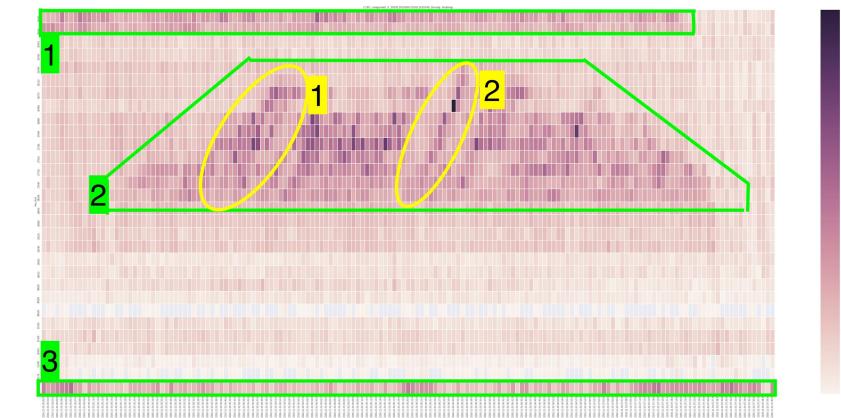
# E18O 2016-11-1 05:00 - 07:40 (component 0)



# E18O 2016-11-1 05:00 - 07:40 (component 0)



*Average speed*



*Density*

# Congestion detection methods

4 congestion detection approaches implemented on batch data

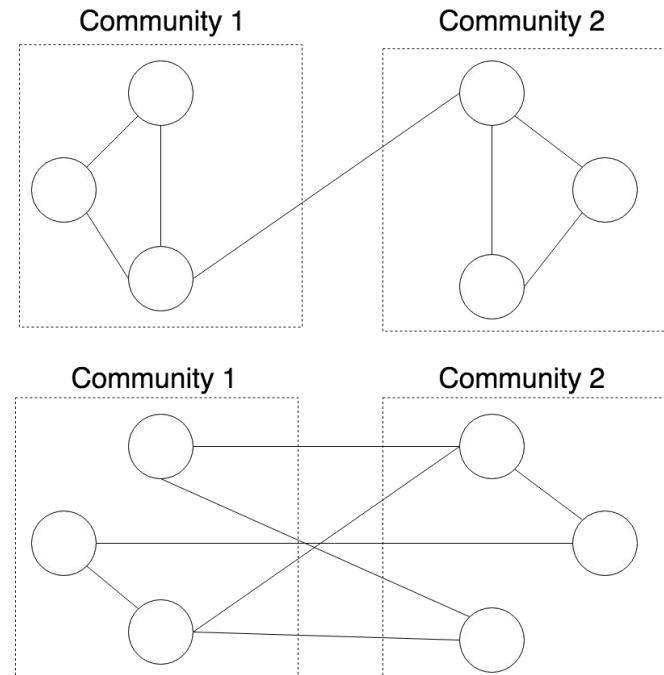
- Louvain modularity
- Hierarchical (data) clustering
- Congested components
- Dengraph

2 best performing implemented as streaming systems in Spark Structured Streaming

- Congested components
- Dengraph

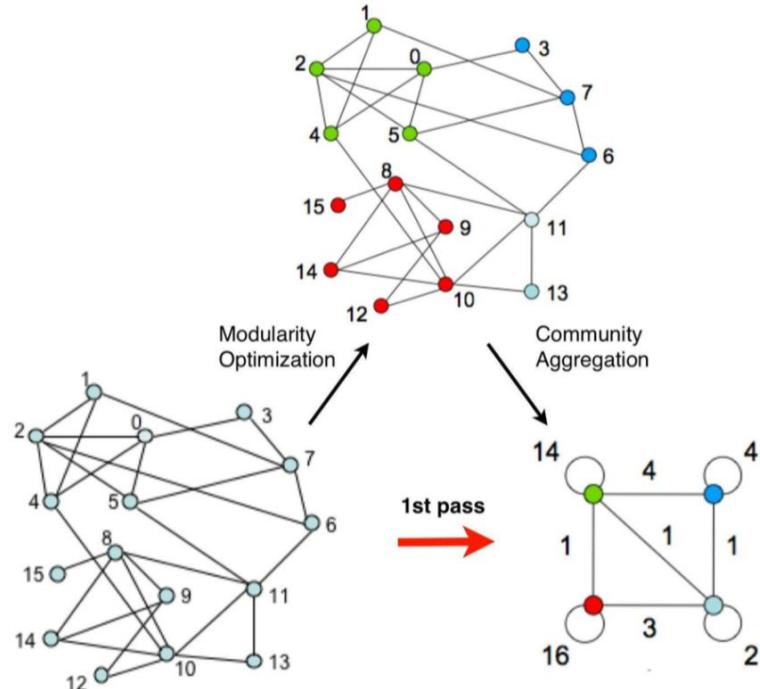
# Louvain modularity

- Weighted modularity optimization community detection algorithm
- Modularity Q
  - Measure of quality of discovered communities, single number in range  $Q = [-1,1]$
  - Fraction of edges in the graph that connect vertices within the same community, minus the expected fraction in a graph with the same community division and degree distribution, but edges connected randomly
  - How much better is the community division than would happen randomly
  - $Q$  approaches 1  $\rightarrow$  good community division
- Idea: Find communities of congested sensors

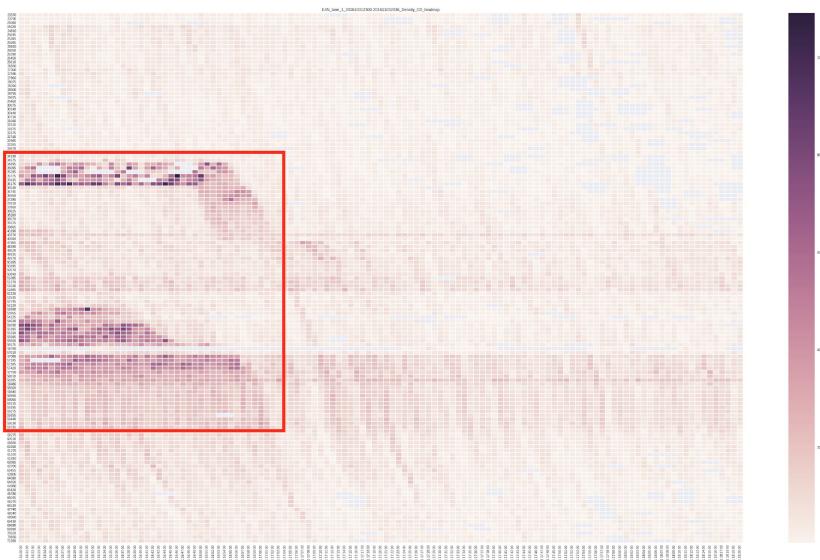


# Louvain modularity

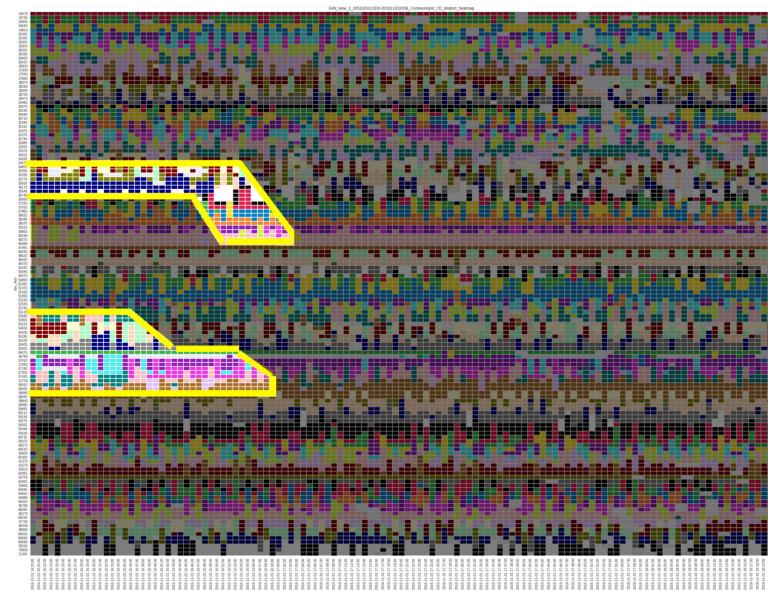
- Edges of graph weighted with density (cars/km)
- Run minute-by-minute on data from E4N, lane 1, 2016-11-1 16:20-18:20
- Average modularity value achieved (averaged over all minutes)
  - $Q = 0.728$  (high)
  - Finishes in one iteration
- Poor results
- Comparison to random graph not suitable for chain-like base graph
  - Assumes each node can connect to any other node



# Louvain modularity results



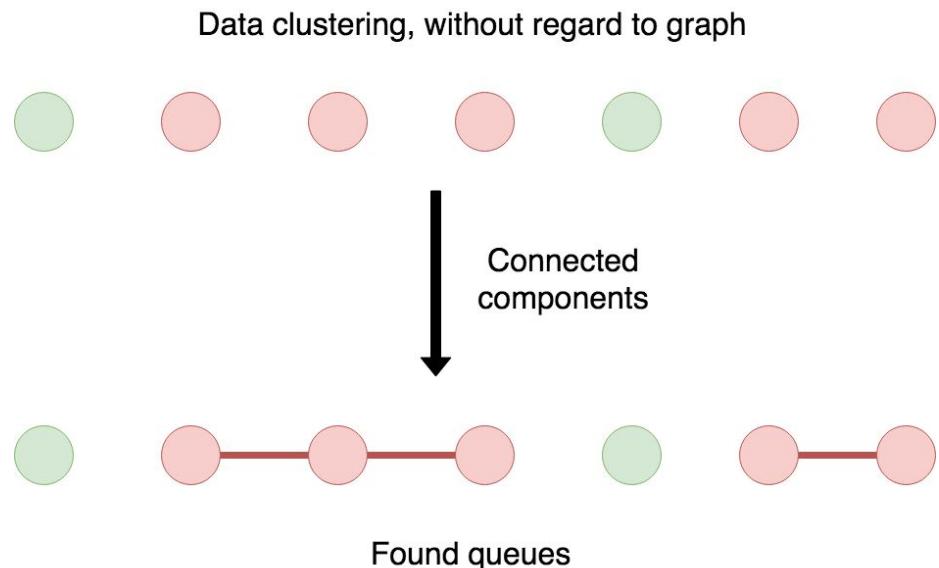
*Observed congestion pattern  
(density)*



*Detected community division*

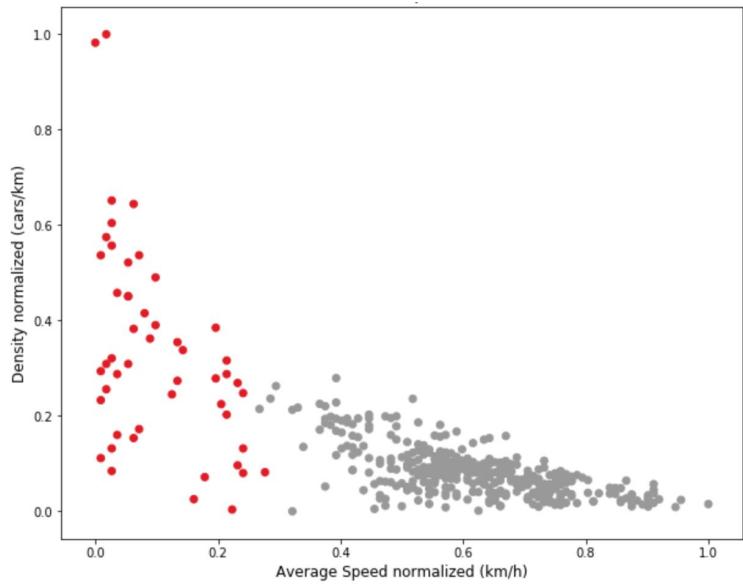
# Hierarchical data clustering

- Base graph is very simple (chain of vertices)
- Idea:
  - Data clustering to find congested sensors (i.e. not graph)
  - Run connected components on graph to find connected component of congested sensors (queue)
- Hierarchical clustering chosen to explore clustering result at different hierarchies
  - Congestion hierarchy



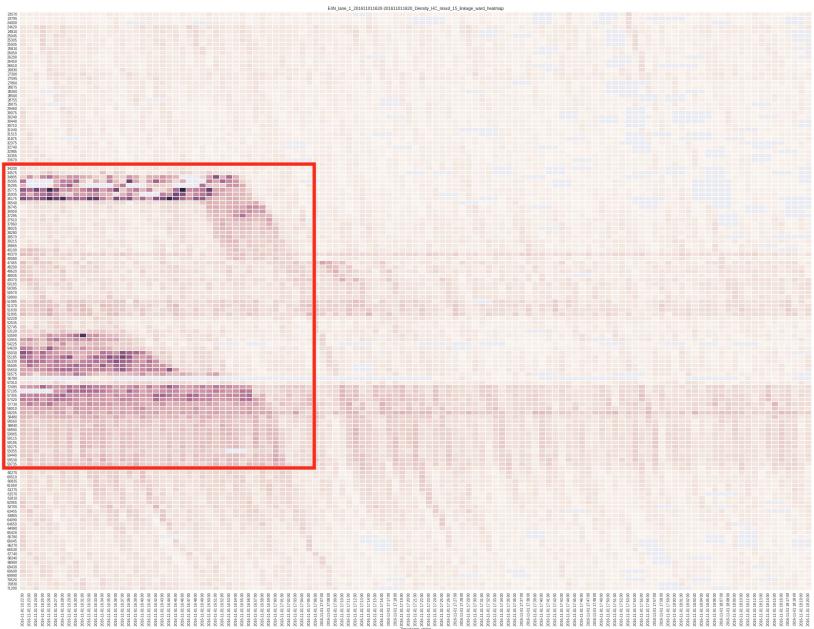
# Hierarchical data clustering

- Features:
  - Average speed and density
  - Normalized for each minute
- Run minute-by-minute on data from E4N, lane 1, 2016-11-1  
16:20-18:20
- Dendrogram cut to generate two clusters
  - Free flow and congested
- Ward's variance minimization linkage criteria gave best results



*Clustering result for 2016-11-01 16:20.  
Ward linkage. 2 clusters.*

# Hierarchical data clustering - result

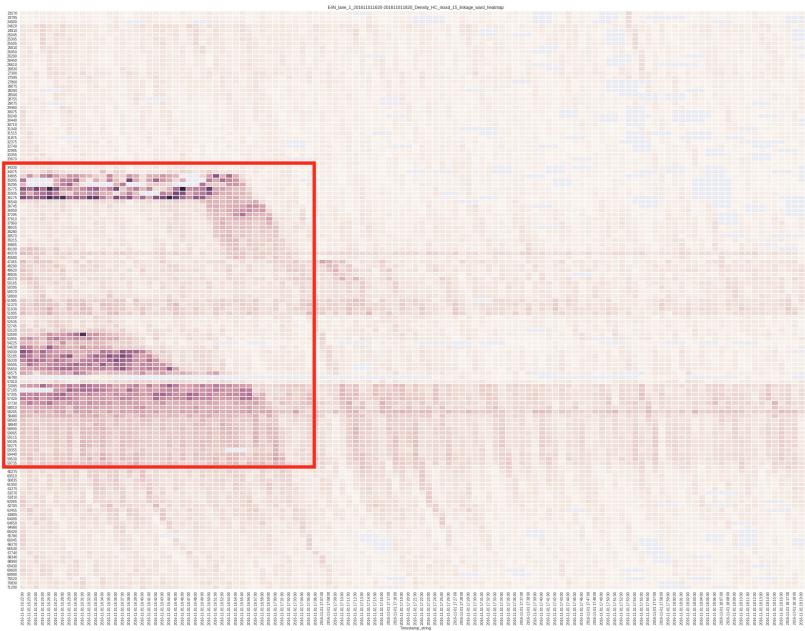


*Observed congestion pattern  
(density)*

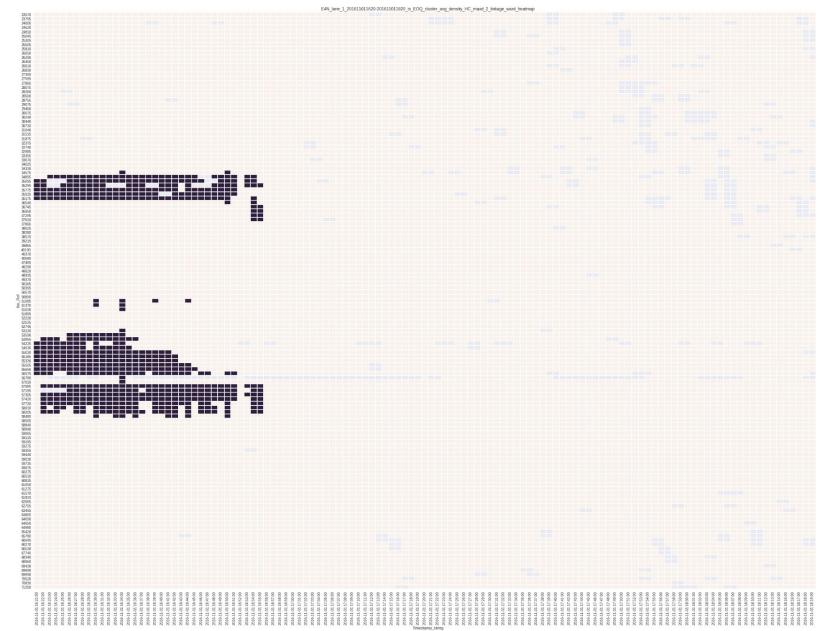


*Cluster division (before split)*

# Hierarchical data clustering - results



Observed congestion pattern  
(density)

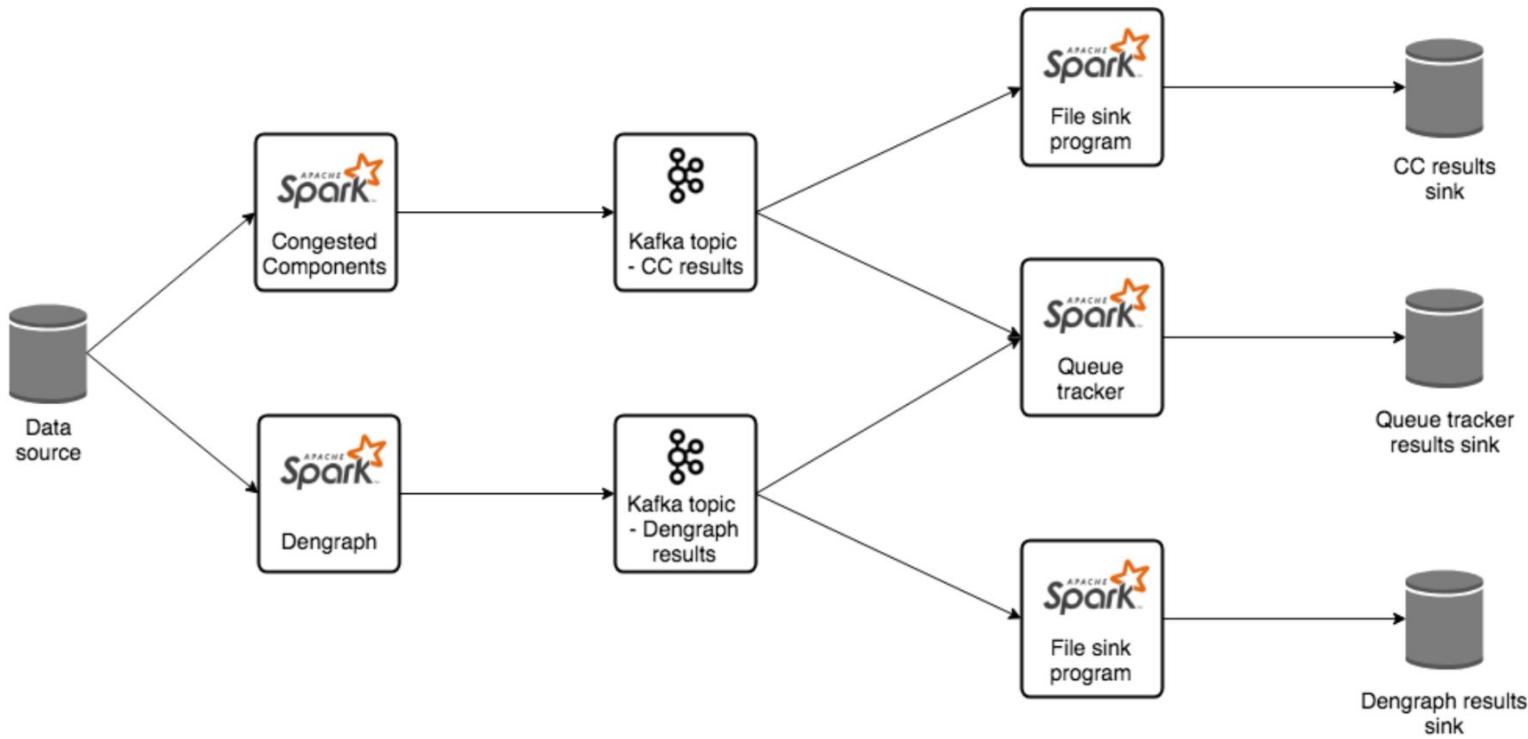


Queue clusters - after split  
Queue detection heuristic:  
avg density within cluster > 30

# Streaming approaches

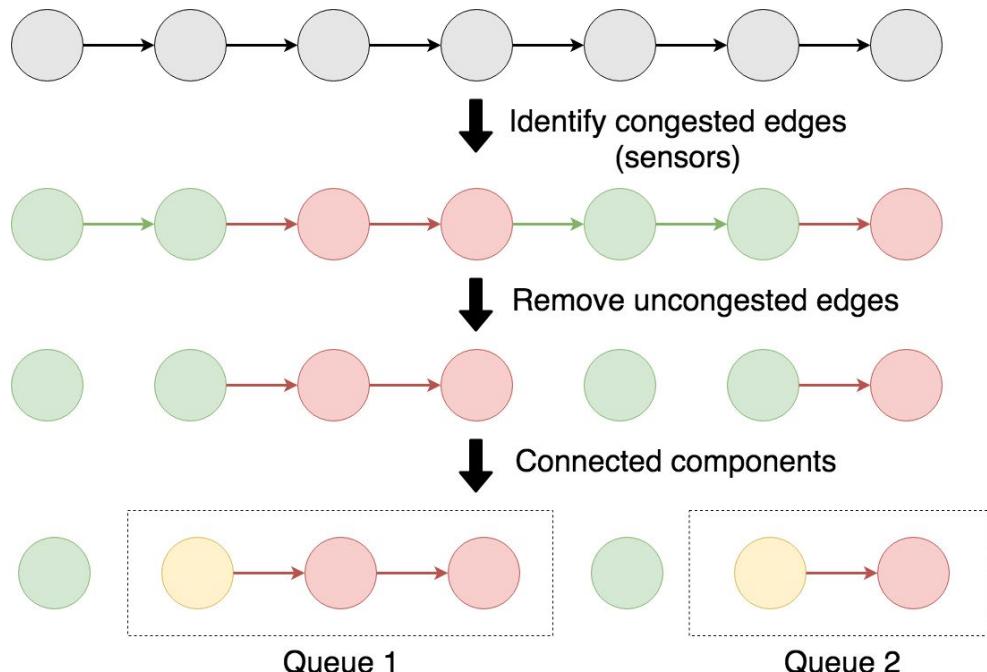
- Congestion detection
  - Congested components
  - Dengraph
- Queue tracker

# System architecture



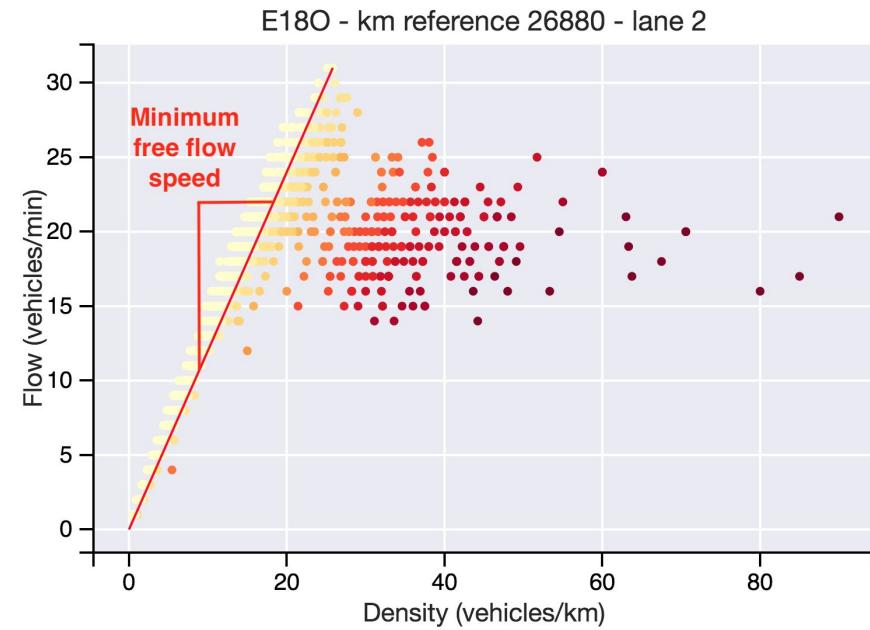
# Congested components

- Idea: Find connected components of congested sensors
- They represent queues



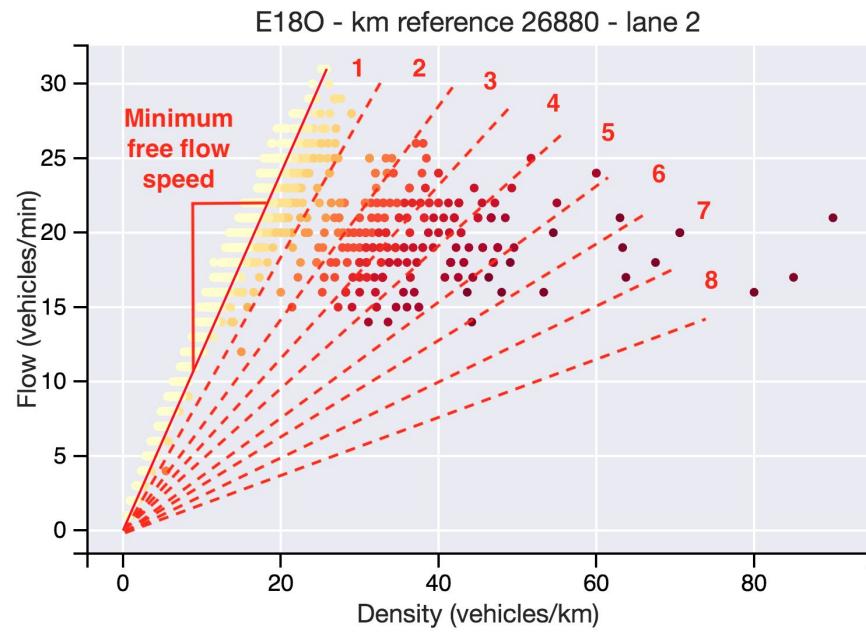
# Congested components

- Detect congestion based on speed
- Congestion threshold
  - Minimum free flow speed
  - Slope of line from origin to maximum free flow point
  - Determined for each individual sensor from historical data
  - Variable between sensors
- Congestion classes
  - Measure of severity of congestion
  - Classify sensor measurements based on how far below minimum free flow speed
    - (increments of 5 km/h)

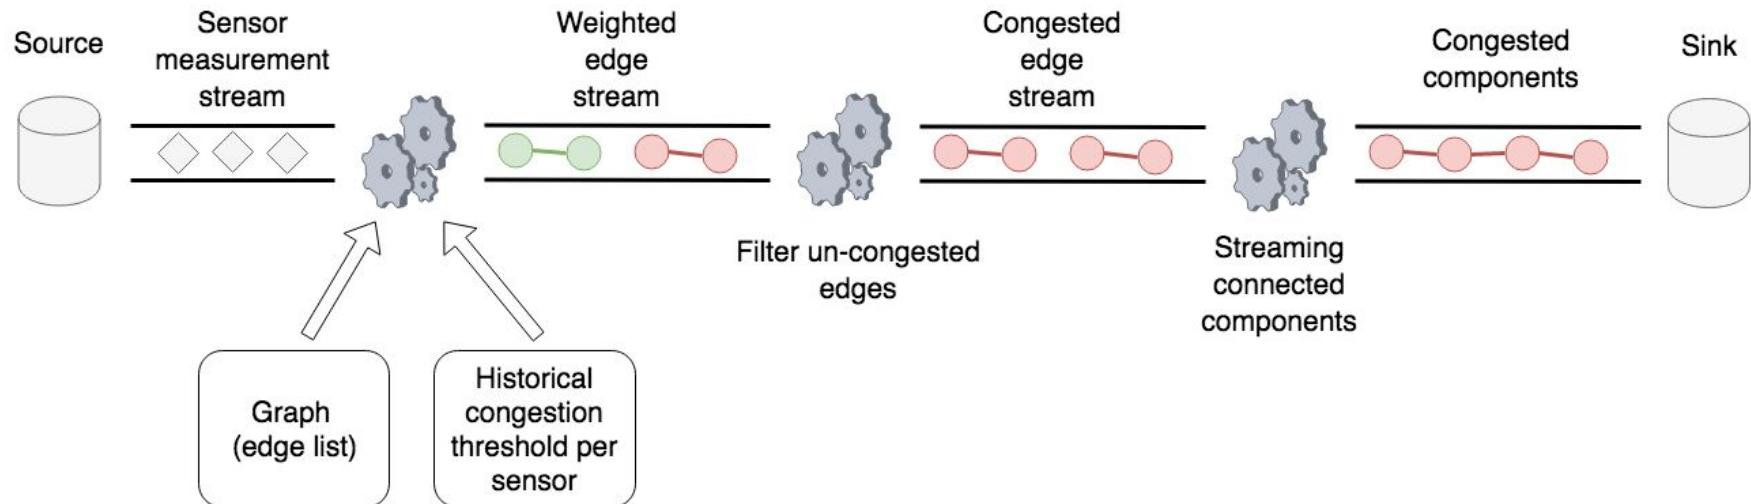


# Congested components

- Detect congestion based on speed
- Congestion threshold
  - Minimum free flow speed
  - Slope of line from origin to maximum free flow point
  - Determined for each individual sensor from historical data
  - Variable between sensors
- Congestion classes
  - Measure of severity of congestion
  - Classify sensor measurements based on how far below minimum free flow speed
    - (increments of 5 km/h)

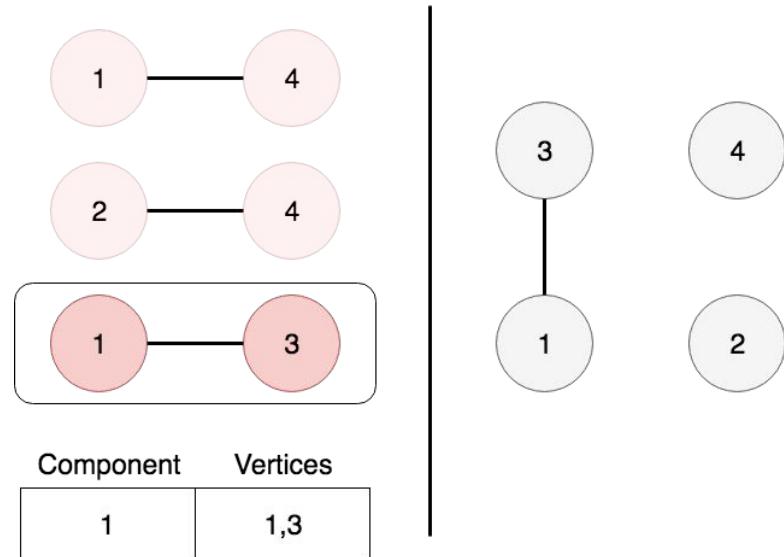


# Congested components



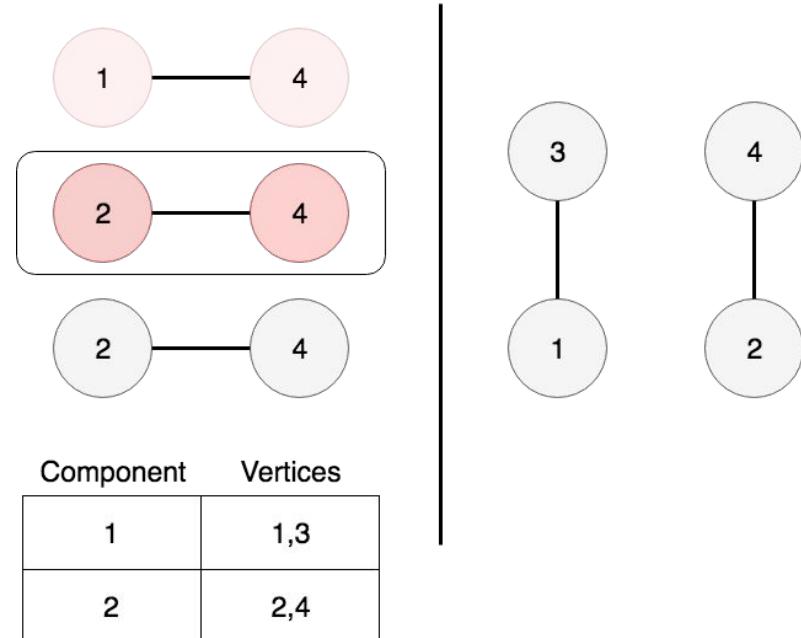
# Streaming connected components

- Single-pass algorithm
- For each edge  $(u,v)$ 
  - If neither  $u$  nor  $v$  in component, create component with ID  $\min(u.ID, v.ID)$
  - If  $u, v$  in different components  $C_1$  and  $C_2$ , merge components and update component ID to  $\min(C_1.ID, C_2.ID)$
  - If only  $u$  or  $v$  in component, add the other one to the component
- State:  
Component ID to vertex IDs map
- State built up for each minute
  - One minute tumbling window



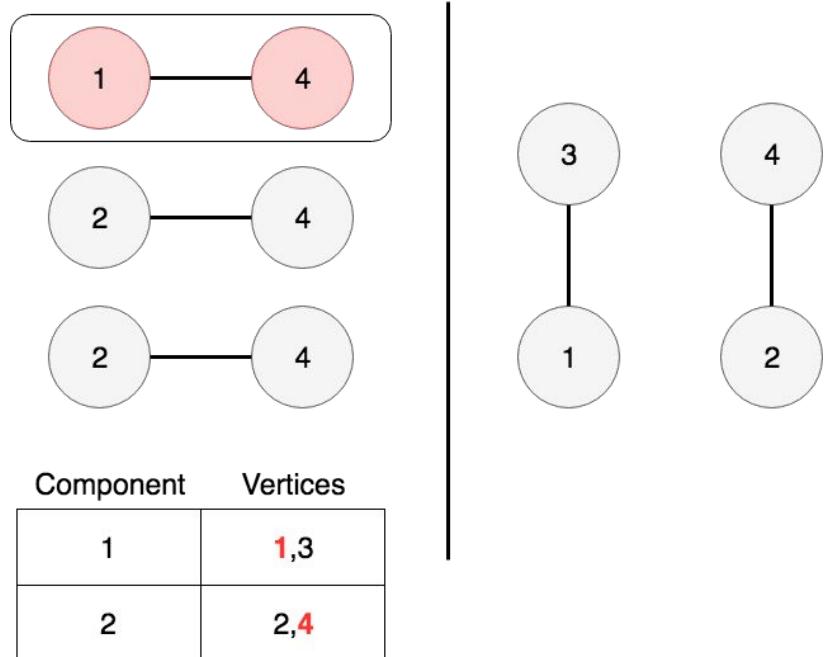
# Streaming connected components

- Single-pass algorithm
- For each edge  $(u,v)$ 
  - If neither  $u$  nor  $v$  in component, create component with ID  $\min(u.ID, v.ID)$
  - If  $u, v$  in different components  $C_1$  and  $C_2$ , merge components and update component ID to  $\min(C_1.ID, C_2.ID)$
  - If only  $u$  or  $v$  in component, add the other one to the component
- State:  
Component ID to vertex IDs map
- State built up for each minute
  - One minute tumbling window



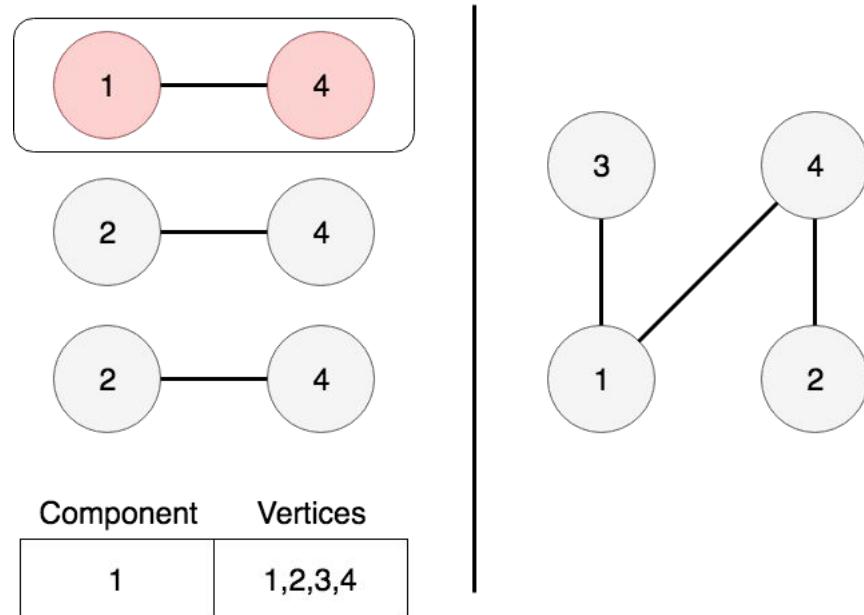
# Streaming connected components

- Single-pass algorithm
- For each edge  $(u,v)$ 
  - If neither  $u$  nor  $v$  in component, create component with ID  $\min(u.ID, v.ID)$
  - If  $u, v$  in different components  $C_1$  and  $C_2$ , merge components and update component ID to  $\min(C_1.ID, C_2.ID)$
  - If only  $u$  or  $v$  in component, add the other one to the component
- State:  
Component ID to vertex IDs map
- State built up for each minute
  - One minute tumbling window



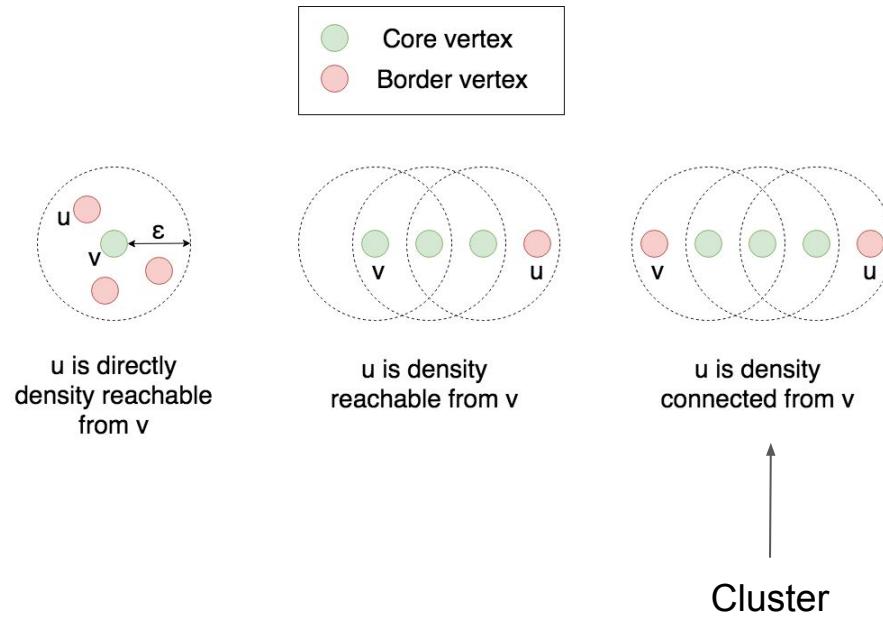
# Streaming connected components

- Single-pass algorithm
- For each edge  $(u,v)$ 
  - If neither  $u$  nor  $v$  in component, create component with ID  $\min(u.ID, v.ID)$
  - If  $u, v$  in different components  $C_1$  and  $C_2$ , merge components and update component ID to  $\min(C_1.ID, C_2.ID)$
  - If only  $u$  or  $v$  in component, add the other one to the component
- State:  
Component ID to vertex IDs map
- State built up for each minute
  - One minute tumbling window



# Dengraph

- Density based graph clustering
- Parameters
  - $\epsilon$  - neighborhood radius (distance)
  - $\eta$  - min number of vertices within neighborhood for a vertex to be a core vertex
- Classify vertices
  - Core, border, noise
- Cluster:
  - Maximal set of vertices that are density connected w.r.t.  $\epsilon, \eta$
  - Nodes not in cluster are classified as noise



# Dengraph distance function

- Distance function

$$dist(p, q) = \begin{cases} 0 & p = q \\ \frac{1}{W_{p,q}} & \exists(p, q) \in E \\ \text{undefined} & \text{otherwise,} \end{cases}$$

- Edges weighted with density
- Essentially thresholds on density
- Critical density about the same for all roads [Highway Capacity Manual 2010]
  - $\approx 28$  vehicles/km/lane

Neighborhood radius $\epsilon$	Corresponding density (vehicles/km)
0,025	40
0,03	33,3
0,035	28,6
0,04	25
0,045	22,2
0,05	20

# Dengraph

- Incremental clustering algorithm
  - Weighted edge stream
- Each edge can have a *positive* or *negative* effect on clustering
  - Positive if edge not seen before or it brings two vertices close enough to fall in each others neighborhoods
    - Cluster creation, absorption or merge
  - Negative if edge removed or if distance between two vertices grows so they are no longer in neighborhood
    - Cluster removed, shrunk or split
- State not carried between minutes
  - Only positive updates (new graph every minute)
  - One minute tumbling window
- State: All edges seen so far, vertex state for all vertices (cluster ID, core /border/noise)

# Dengraph noise resistance

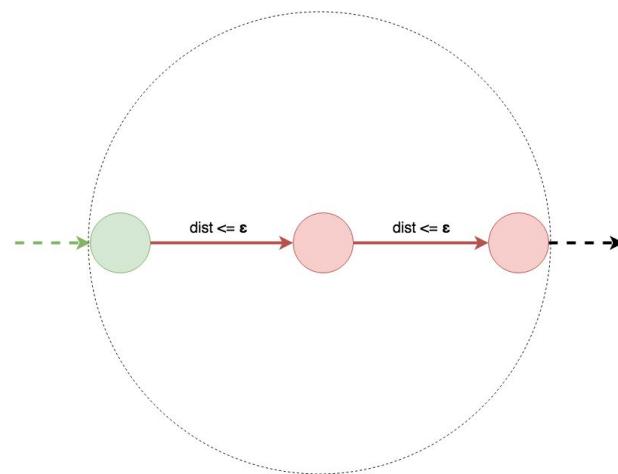
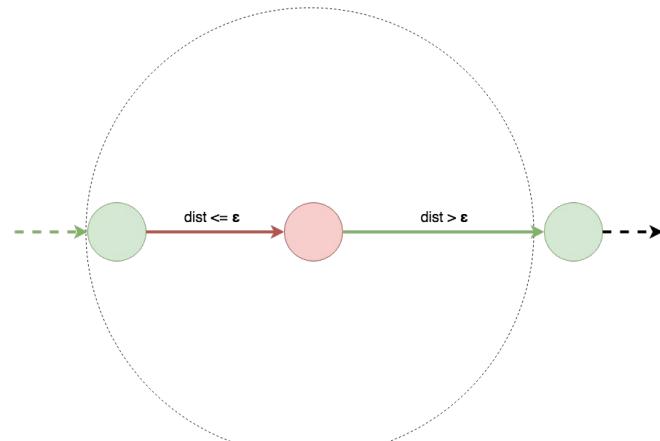
Resistant to noise through  $\eta$  parameter

- Minimum number of nodes in neighborhood to be a core vertex

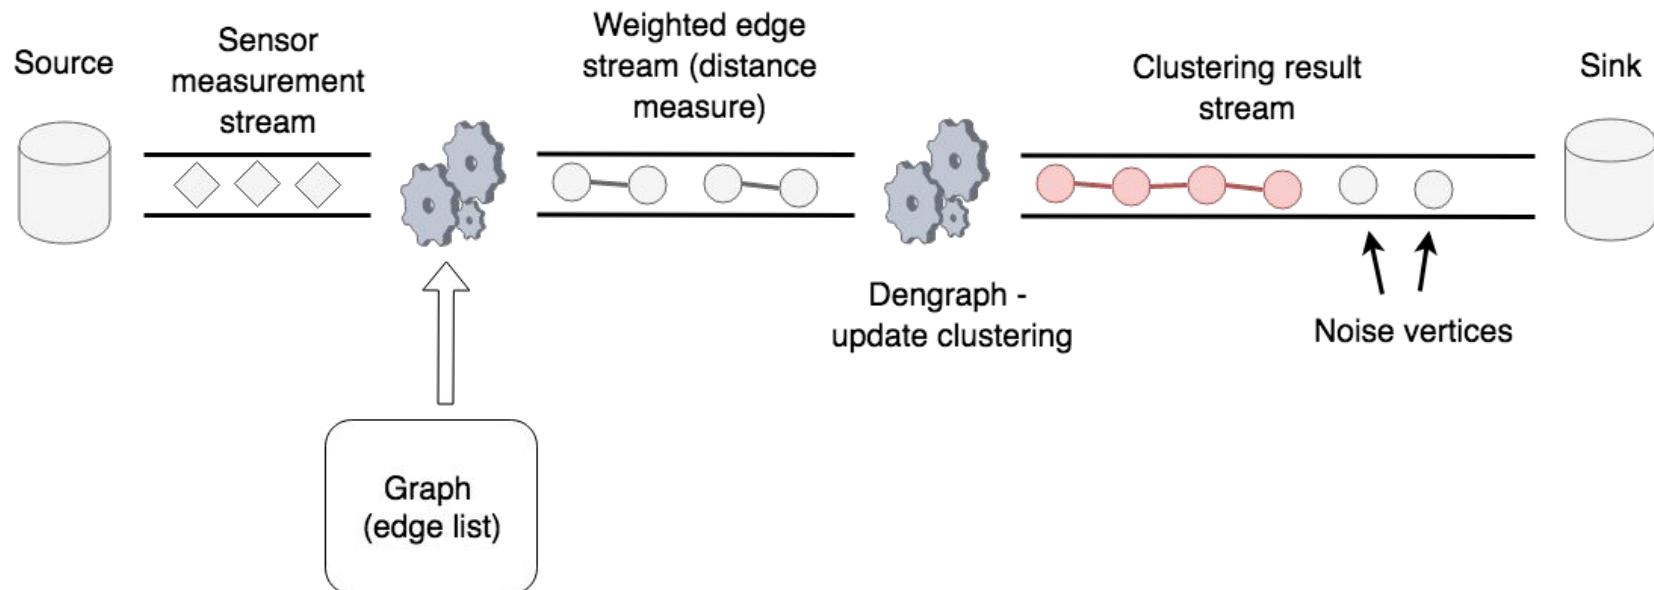
Used  $\eta = 2$

- Need two vertices in neighborhood to form a cluster
  - At least two congested sensors to form a queue (+ a 3rd vertex)

Noise gets classified as noise vertices

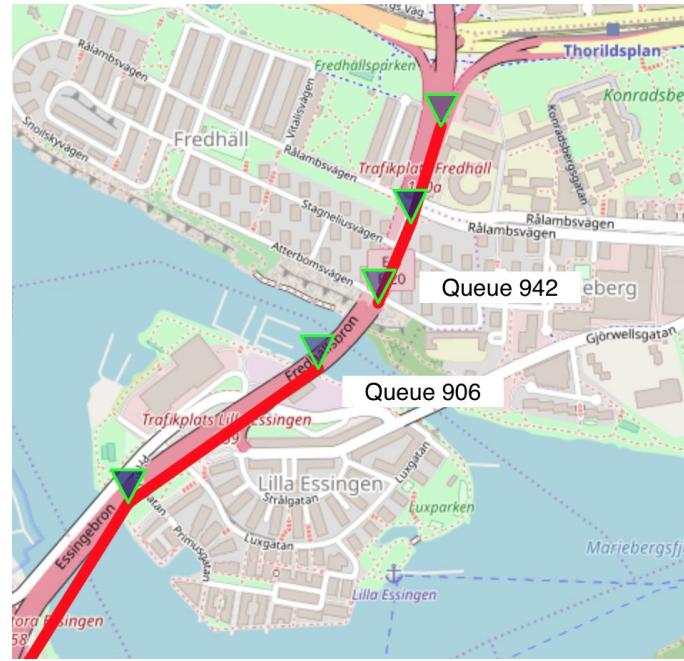
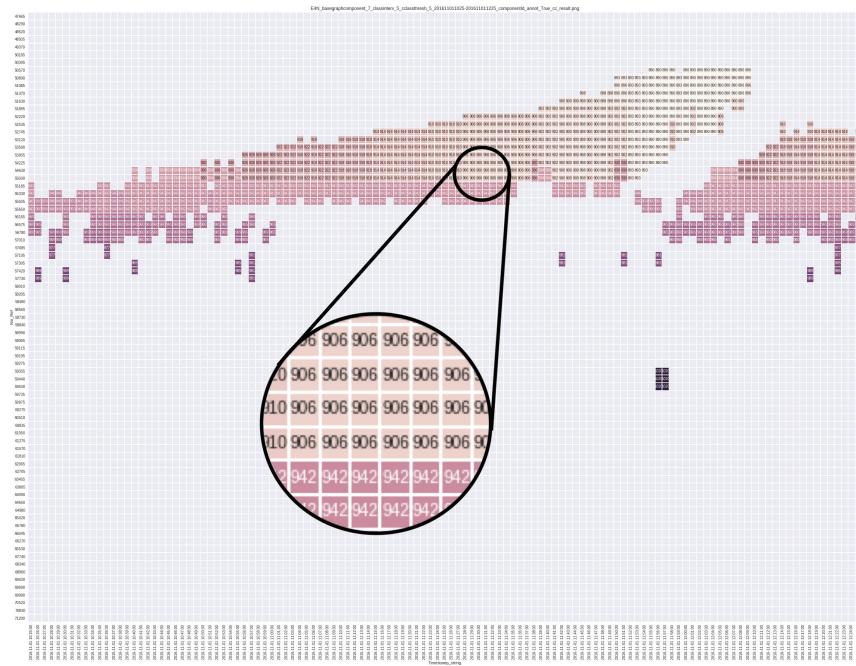


# Dengraph



# Detecting individual queues

Both algorithms find individual queues, with unique queue IDs (within minute)



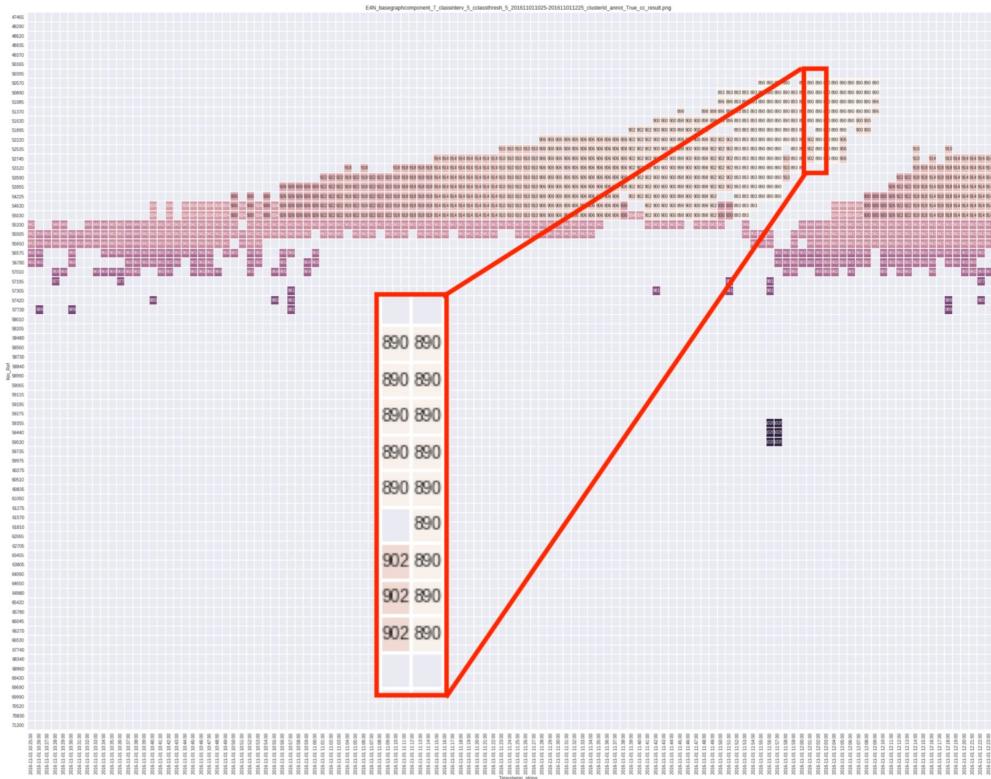
# Queue tracker (through time)

- Implemented as separate streaming system that reads CC/DenGraph results from Kafka
- Calculate Jaccard distance between clusters at time t and clusters at time t-1, pairwise
- Distance threshold 1
  - If Jaccard distance < 1 then we found a parent queue
  - Only one sensor needed in overlap
- Naive, scalability concerns (pairwise comparison -  $O(n^2)$ )
  - Worst case - each sensor in its own cluster (won't happen)
  - 2037 sensors at most in the graph (2016-5-9 onwards)
  - $N(N-1)/2 = 2.073.666$  cluster comparisons
  - CC worst case 518.671 (2 vertices per cluster)
  - DenGraph worst case 230.181 (3 vertices per cluster)
- Run only on *final* results for each pair of minutes

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

$$d_J(A, B) = 1 - J(A, B)$$

# Queue tracker (through time)



# Accuracy evaluation

- Congestion detection algorithms evaluated w.r.t. percentage of labeled queues and shockwaves they detect (recall)
  - As determined by visual inspection
  - Domain expert approach
- Criteria
  - Detected shape conforms to shape in underlying data
    - Emphasis on upstream congestion boundary (EOQ)
  - Boundaries are clear
  - Noise at boundary minimal

	Found queues	Queue recall	Found s.waves	S.wave recall	WSM
CC (c. class 1)	1	6.3%	0	0%	0.05
CC (c. class 3)	5	31.3%	1	6.7%	0.25
CC (c. class 5)	<b>15</b>	<b>93.8%</b>	10	66.7%	<b>0.87</b>
CC (c. class 7)	13	81.3%	12	80%	0.81
CC (c. class 9)	9	56.3%	<b>14</b>	<b>93.3%</b>	0.66
DG ( $\epsilon = 0, 025$ )	3	18.8%	10	66.7%	0.31
DG ( $\epsilon = 0, 03$ )	4	25%	11	73.3%	0.37
DG ( $\epsilon = 0, 035$ )	13	81.3%	11	73.3%	0.79
DG ( $\epsilon = 0, 04$ )	14	87.5%	6	40%	0.76
DG ( $\epsilon = 0, 045$ )	12	75%	2	13.3%	0.60
DG ( $\epsilon = 0, 05$ )	5	31.3%	0	0%	0.23

*Accuracy evaluation.*

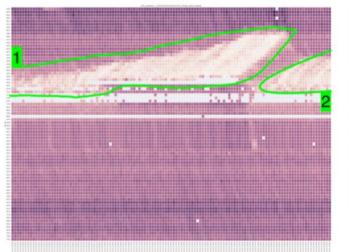
*16 queues, 15 shockwaves in labeled data set.  
WSM 75% weight on queues, 25% weight on  
shockwaves*

# Congested components results

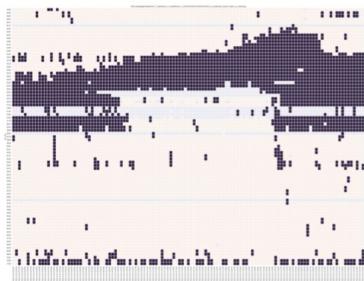
Run with different congestion class thresholds (1-9)

Congestion class threshold:

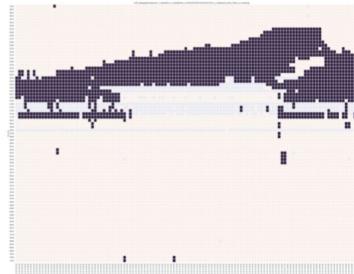
- 1 - No queue visible, noise
- 3 - Queue 1 detected
- 5 - Queue 1 and 2 detected



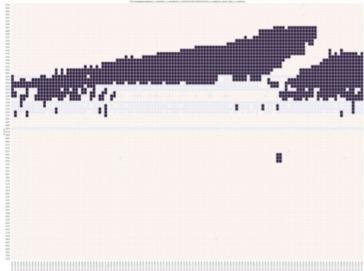
(a) Measured average speed values. Two queues delimited by green borders.



(b) Congestion class threshold 1



(c) Congestion class threshold 3



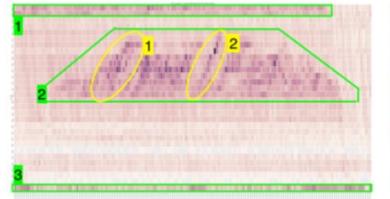
(d) Congestion class threshold 5

*Test congestion pattern 3*

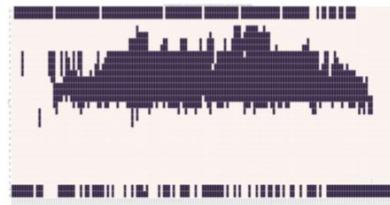
*E4N, component 7, 2016-11-01 10:25-12:25*

# Dengraph results

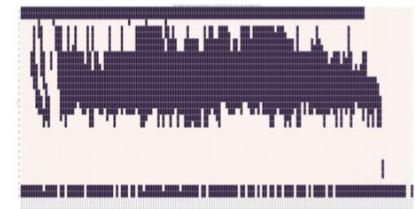
- Run with different  $\epsilon$  values
  - Neighborhood radius (distance)
- $\eta = 2$  constant
  - min number of vertices in neighborhood
- $\epsilon = 0,05$ 
  - Queues 1 and 3 detected, no shockwaves
- $\epsilon = 0,04$ 
  - All queues detected, no shockwaves
- $\epsilon = 0,03$ 
  - Queue 2 detected, both shockwaves detected



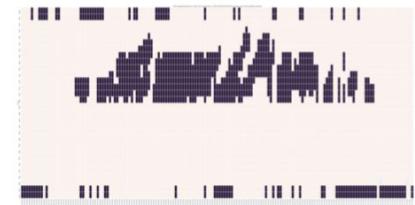
(a) Measured density values. 3 queues delimited in green, 2 shockwaves in yellow.



(c)  $\epsilon = 0,04$



(b)  $\epsilon = 0,05$

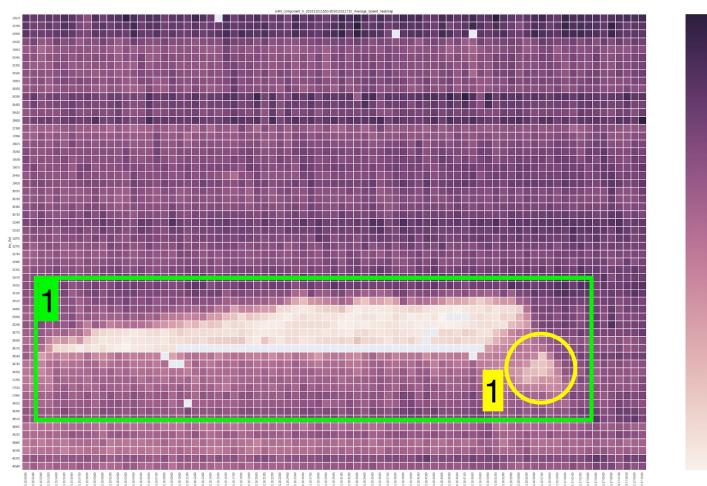


(d)  $\epsilon = 0,03$

*Test congestion pattern 6  
E18O, component 0, 2016-11-01 05:00-07:40*

# Performance evaluation

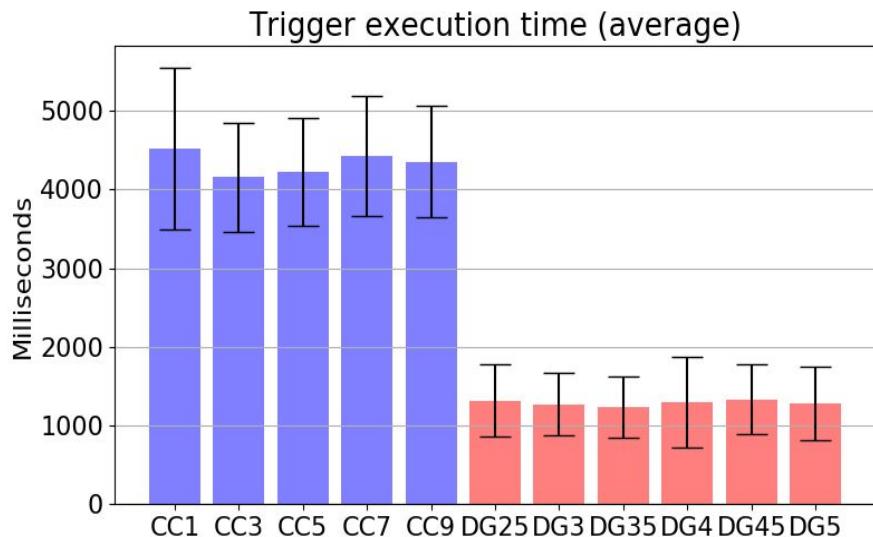
- Algorithms run on data from 2016-11-01 15:50-17:10
- Evaluation w.r.t.
  - Time to process each minute
  - Throughput (rows/sec)
  - Memory footprint (for each minute)
- Sensor measurements from **entire graph** is streamed
  - Average 1837 measurements per minute (after filtering errors)
- Evaluation on local machine
  - 2.5 GHz quad core, 16GB RAM, Spark v2.3.1 (8 threads), Kafka v1.1



Test congestion pattern 1.  
E4N component 0. 2016-11-01  
15:50-17:10.

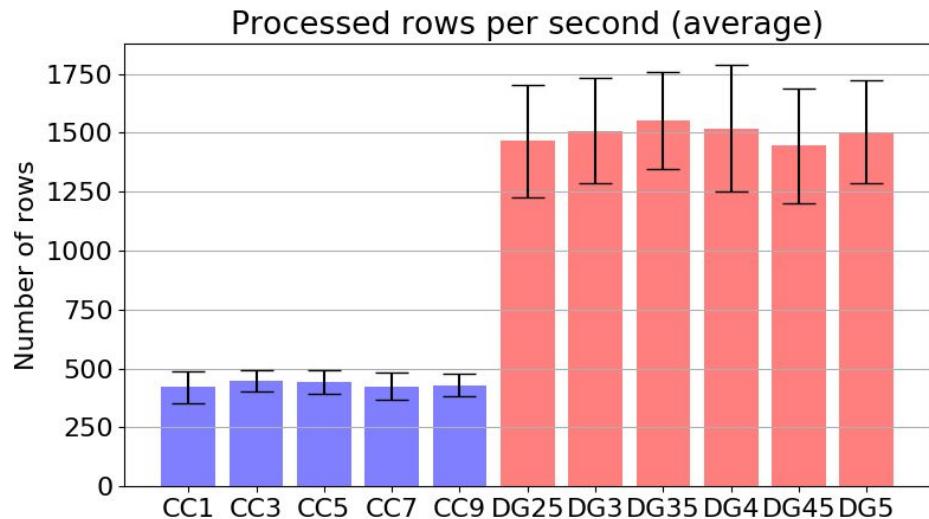
# Trigger execution time

- Stream source set up so that entire minute's worth of data processed in a trigger
- Average values over the 80 minutes of the time span
- Time measures entire pipeline, not just core algorithm
- Dengraph twice as fast



# Processed rows per second

Dengraph more than twice as fast



# Memory usage per trigger

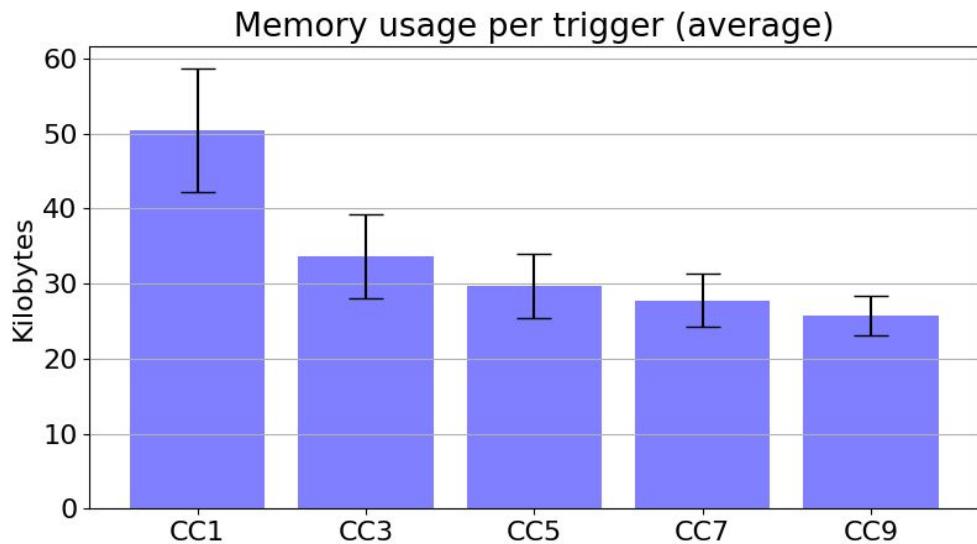
State from two adjacent minutes  
in memory at any given point

Congested components

- Clear trend
- Less memory with higher congestion threshold

Dengraph

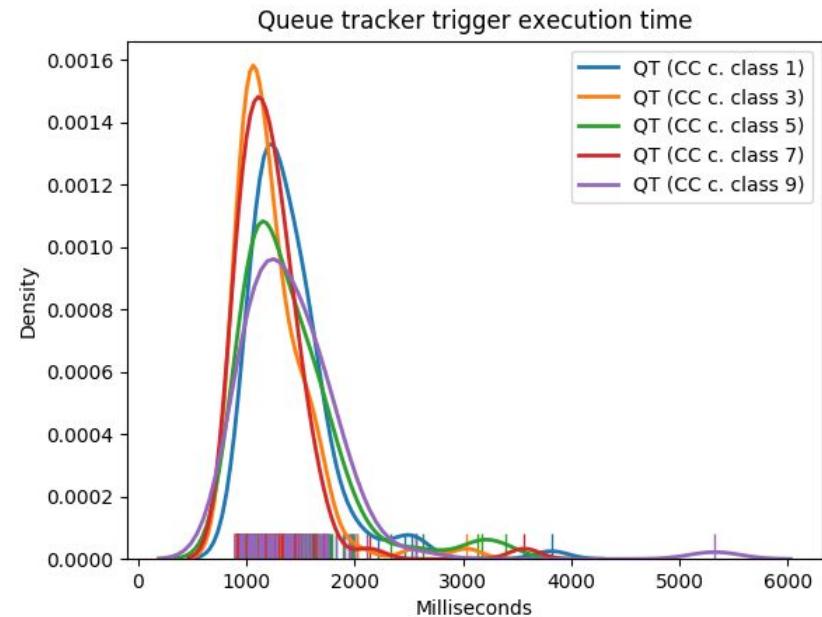
- Constant at 965.71 KB on average



# Queue tracker - Trigger execution time

Uneven work performed in each trigger

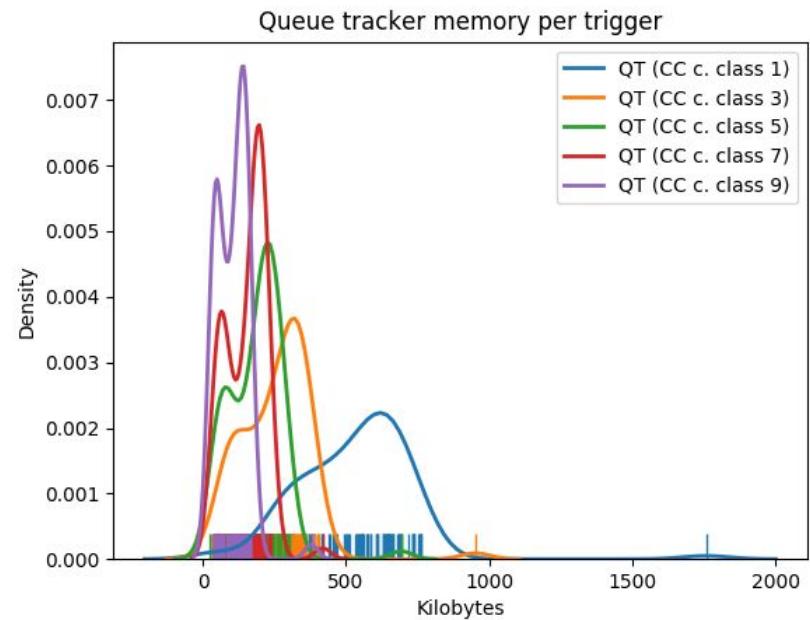
- Only computes parent clusters when all final rows from pair of minutes to be received
- Otherwise, just accumulates final rows in state
- May compute parent clusters for a number of minute-pairs in same trigger



# Queue tracker - Memory per trigger

Clear trend

- Less memory with higher congestion threshold



# Conclusions

- Congested components better accuracy overall
  - Able to detect queues with accuracy at best up to  $\approx 94\%$  and shockwaves  $\approx 93\%$
- Dengraph detects at best  $\approx 87.5\%$  of queues and  $\approx 73.3\%$  of shockwaves
- Dengraph however more than twice as fast
  - $\approx 1.5$  seconds to process each minute while taking up more space
- Tradeoff between detection of queues (congestion areas) and shockwaves (detail in congestion areas)
- Louvain modularity not suitable for simple road network graph
- Hierarchical data clustering approach feasible
  - Two phases needed
    - Clustering (incremental)
    - Connected components

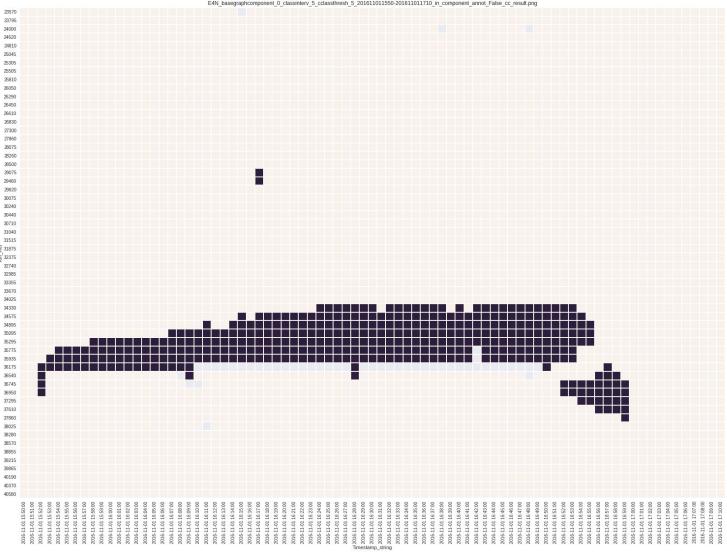
# Future work

- Scalability considerations
  - Evaluation and graph partitioning
- Ground truth and accuracy evaluation strategy
  - Traffic simulators
  - Per sensor congestion labelling by human expert
- End-of-queue warning system
  - Use reachability graph to send warnings to sensors/cars approaching end-of-queue
- Dedicated end-of-queue detection system
  - Detect high risk jumps in density or drops in average speed
- Integrate with streaming predictions
  - Predict congestion, existence and movement
- Evolutionary clustering
- Continuously update minimum free flow speed per sensor

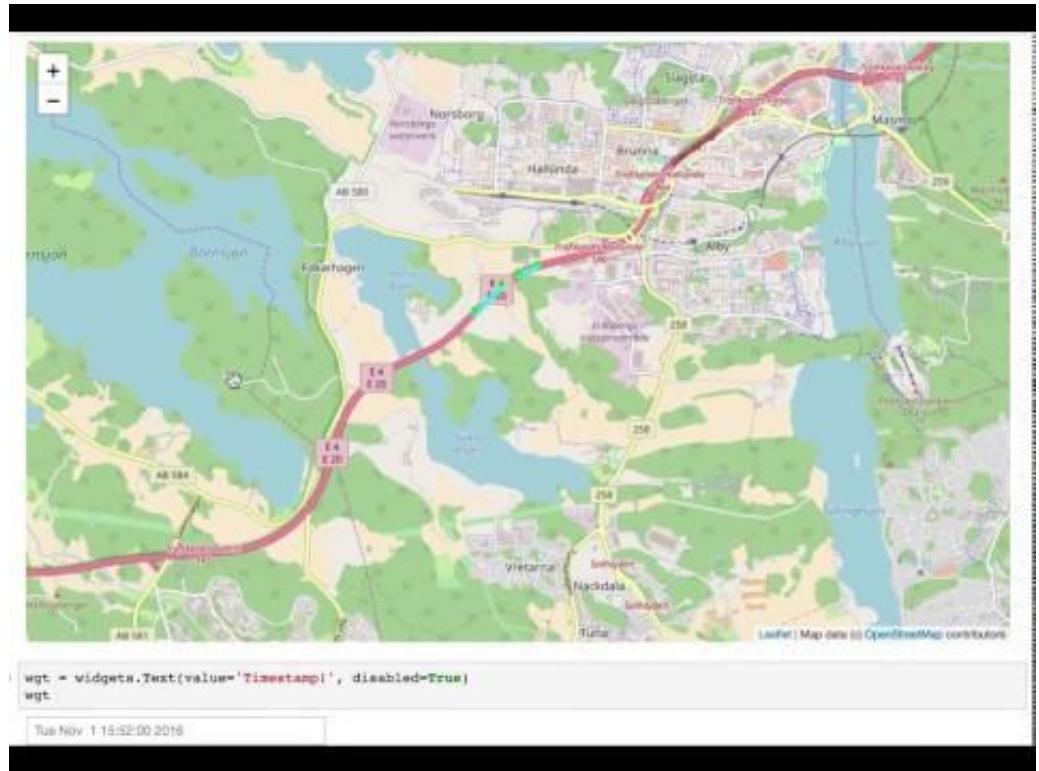
Thank you!

# Appendix

# Video visualization



*E4N, component, 0, 2016-1101  
15:50-17:10  
Congested components, cclass 5*



# Louvain modularity - procedure

2 phases

Phase 1

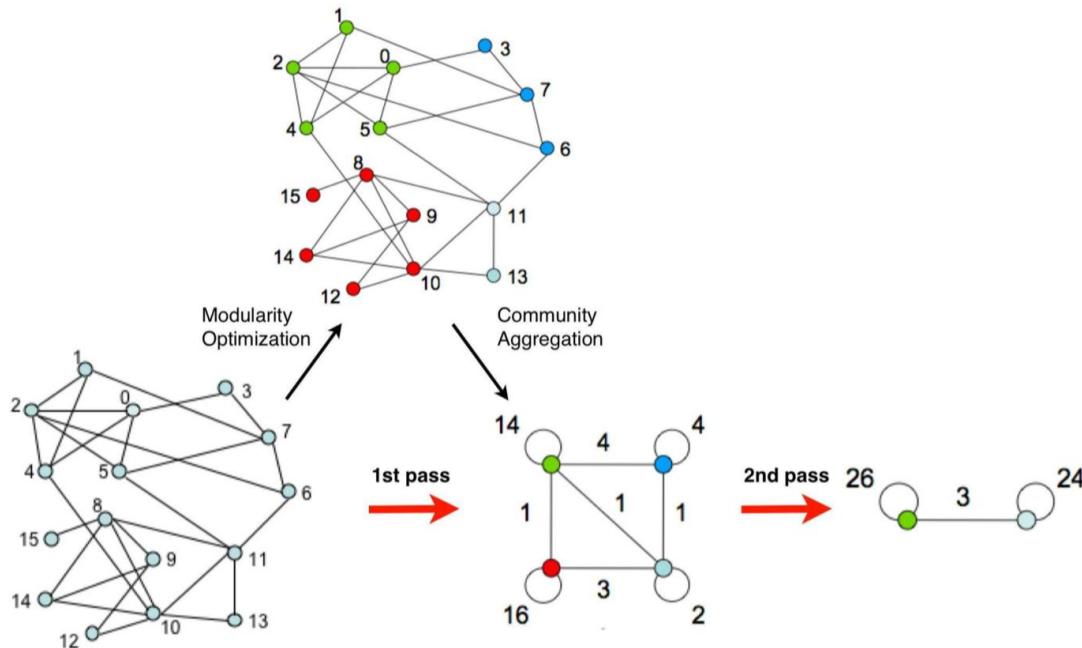
- For each node - move into community of neighbor which gives the greatest gain in modularity
- Repeat until no further gain in modularity can be achieved

Phase 2

- New graph constructed
- Communities found in phase 1 are the vertices
- The weights of the edges between the community-vertices in this new graph are the sum of the weights of edges between vertices in the two communities that the community-vertices we are linking represent

Repeat until no changes made in phase 1 and modularity has been maximized

# Louvain modularity



Blondel et al.

# Weighted modularity

$$Q = \frac{1}{2m} \sum_{i,j} \left[ A_{ij} - \frac{k_i k_j}{2m} \right] \delta(c_i, c_j),$$

$A_{\{i,j\}}$  := weight of edge connecting  $i, j$

$k_i$  = sum of all weights connecting to  $i$

$C_i$  = community of  $i$

$\delta$  = 1 if both same, 0 otherwise

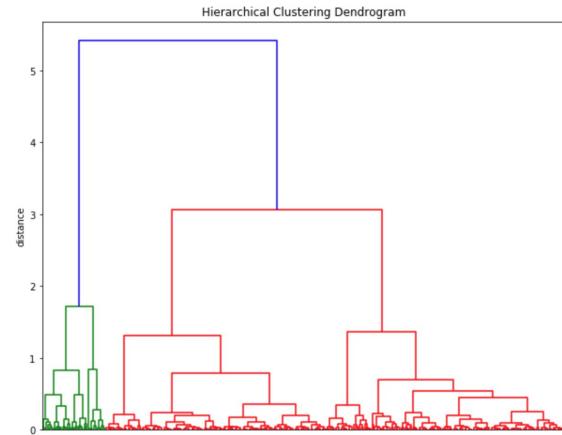
$m = \frac{1}{2} * \text{sum}_{\{i,j\}} A_{\{i,j\}}$  i.e. sum of all weights / half

(Sum of all weights within communities - (sum of all weights to  $i$  \* sum of all weights to  $j$ )/sum of all weights)  
/ sum of all weights

# Hierarchical clustering Ward's variance minimization

## Hierarchical clustering

- Start with all in separate clusters
- Merge clusters closer to each other depending on distance metric - stepwise
- Builds hierarchy (dendrogram)



Ward's variance minimization distance function:

$$d(u, v) = \sqrt{\frac{|v| + |s|}{T} d(v, s)^2 + \frac{|v| + |t|}{T} d(v, t)^2 - \frac{|v|}{T} d(s, t)^2}$$

where  $u$  is the newly joined cluster consisting of clusters  $s$  and  $t$ ,  $v$  is an unused cluster in the forest,  $T = |v| + |s| + |t|$ , and  $| * |$  is the cardinality of its argument. This is also known as the incremental algorithm.

# Dengraph positive updates

Positive updates can lead to

- Cluster creation
  - New core vertex
  - Create new cluster
- Cluster absorption
  - Noise vertices now in core vertex neighborhood
  - Noise vertices absorbed to core vertex's cluster
- Cluster merge
  - Core vertices from different clusters now in each others neighborhoods
  - Merge clusters

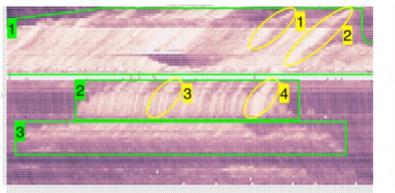
# Congested components results

Run with different congestion class thresholds (1-9)

Accuracy evaluation:

Congestion class threshold

- 5 - All 3 queues, shockwaves 2,3,4
- 7 - Queue 1,2 detected, all shockwaves detected
- 9 - Queue 1 detected, shockwaves 1,2,4 clear



(a) Measured average speed values. 3 queues delimited in green, 4 shockwaves in yellow.



(b) Congestion class threshold 5



(c) Congestion class threshold 7



(d) Congestion class threshold 9

*Test congestion pattern 4  
E4N, component 7, 2016-11-01 13:15-16:15*

# The traffic sensor graph

## Time-span graphs

- Sensors added and removed throughout the years
- Time-span graphs to accurately represent the road system at any point
- Latest graph the largest, 2037 sensors and 2077 edges

<b>Graph no.</b>	<b>Valid from</b>	<b>Valid to</b>	<b>No. of sensors</b>
1	2001-1-1	2015-2-25	1843
2	2015-2-26	2015-11-10	1845
3	2015-11-11	2016-3-26	1917
4	2016-3-27	2016-3-28	1994
5	2016-3-29	2016-5-8	1976
6	2016-5-9	onwards	2037